

project_3

October 8, 2024

Instituto Tecnológico de Aeronáutica - ITA

Divisão de Engenharia Eletrônica - IEE

ET-287 - Processamento de sinais usando redes neurais

Professora Sarah Negreiros de Carvalho Leite

Aluno Felipe Keller Baltor

1 Projeto 3

1.0.1 1. *Baixe a base de dados, disponível em:*
<https://www.kaggle.com/datasets/uciml/faulty-steel-plates?resource=download>.

```
[1]: %reset
```

Once deleted, variables cannot be recovered. Proceed (y/[n])? y

```
[2]: # Importações básicas

import numpy as np # biblioteca de manipulação vetorial e numérica
import matplotlib.pyplot as plt # biblioteca para traçar gráficos
import pandas as pd # biblioteca de manipulação de dados tabulares
from pathlib import Path # biblioteca para manipulação de "paths"
import urllib3 # biblioteca para download do dataset
from tqdm import tqdm # barra de download
import zipfile
import scipy
import polars as pl
```

```
[3]: # Checando se estamos no diretório correto

project_dir = Path('.')

assert project_dir.resolve().name == 'project_3'
```

```

[4]: # Baixando os dados

data_dir = project_dir / 'data'
data_file = data_dir / 'faults.csv'

if not data_file.is_file():
    data_dir.mkdir(exist_ok = True)

    data_compressed = project_dir / 'data.zip'
    if not data_compressed.is_file():
        data_url = "https://storage.googleapis.com/kaggle-data-sets/2363/3972/
↳bundle/archive.zip?
↳X-Goog-Algorithm=GOOG4-RSA-SHA256&X-Goog-Credential=gcp-kaggle-com%40kaggle-161607.
↳iam.gserviceaccount.
↳com%2F20241002%2Fauto%2Fstorage%2Fgoog4_request&X-Goog-Date=20241002T200249Z&X-Goog-Expires

        http = urllib3.PoolManager()
        CHUNK_SIZE = 2**16

        resp = http.request('GET', data_url, preload_content = False)
        TOTAL_SIZE = int(resp.headers.get('Content-Length'))
        with (
            open(data_compressed, 'wb') as file,
            tqdm(
                total = TOTAL_SIZE,
                desc = f'Downloading {data_compressed.name}',
                unit = 'B',
                unit_scale = True) as bar
        ):
            for chunk in resp.stream(CHUNK_SIZE):
                size = file.write(chunk)
                bar.update(size)
            resp.release_conn()

        print("Extracting zip to data folder")
        with zipfile.ZipFile(data_compressed, 'r') as zip_ref:
            zip_ref.extractall(data_dir)
        print("Data is ready!")

    else:
        print("Data already exists")

```

Data already exists

```
[5]: faults = pl.read_csv(data_file)
```

```
[6]: faults.columns
```

```
[6]: ['X_Minimum',
      'X_Maximum',
      'Y_Minimum',
      'Y_Maximum',
      'Pixels_Areas',
      'X_Perimeter',
      'Y_Perimeter',
      'Sum_of_Luminosity',
      'Minimum_of_Luminosity',
      'Maximum_of_Luminosity',
      'Length_of_Conveyer',
      'TypeOfSteel_A300',
      'TypeOfSteel_A400',
      'Steel_Plate_Thickness',
      'Edges_Index',
      'Empty_Index',
      'Square_Index',
      'Outside_X_Index',
      'Edges_X_Index',
      'Edges_Y_Index',
      'Outside_Global_Index',
      'LogOfAreas',
      'Log_X_Index',
      'Log_Y_Index',
      'Orientation_Index',
      'Luminosity_Index',
      'SigmoidOfAreas',
      'Pastry',
      'Z_Scratch',
      'K_Scatch',
      'Stains',
      'Dirtiness',
      'Bumps',
      'Other_Faults']
```

Existe um erro de ortografia na coluna *K_Scatch*, que deveria ser *K_Scratch*, vamos corrigir isso:

```
[7]: faults = faults.rename({"K_Scatch": "K_Scratch"}, strict = False)

faults
```

```
[7]: shape: (1_941, 34)
```

X_Minimum	X_Maximum	Y_Minimum	Y_Maximum	...	Stains	Dirtiness	Bumps
Other_Faults							
---	---	---	---		---	---	---

```

---
i64      i64      i64      i64      i64      i64      i64
i64

42      50      270900    270944    ... 0      0      0
0
645     651     2538079    2538108    ... 0      0      0
0
829     835     1553913    1553931    ... 0      0      0
0
853     860     369370     369415     ... 0      0      0
0
1289    1306    498078     498335     ... 0      0      0
0
...      ...      ...      ...      ... ...      ...      ...
...
249     277     325780     325796     ... 0      0      0
1
144     175     340581     340598     ... 0      0      0
1
145     174     386779     386794     ... 0      0      0
1
137     170     422497     422528     ... 0      0      0
1
1261    1281    87951      87967      ... 0      0      0
1

```

Printando os nomes das colunas, temos:

```
[8]: for index, column in enumerate(faults.columns):
      print(f"Index: {index + 1}, Column Name: {column}")
```

```

Index: 1, Column Name: X_Minimum
Index: 2, Column Name: X_Maximum
Index: 3, Column Name: Y_Minimum
Index: 4, Column Name: Y_Maximum
Index: 5, Column Name: Pixels_Areas
Index: 6, Column Name: X_Perimeter
Index: 7, Column Name: Y_Perimeter
Index: 8, Column Name: Sum_of_Luminosity
Index: 9, Column Name: Minimum_of_Luminosity
Index: 10, Column Name: Maximum_of_Luminosity
Index: 11, Column Name: Length_of_Conveyer
Index: 12, Column Name: TypeOfSteel_A300
Index: 13, Column Name: TypeOfSteel_A400

```

```

Index: 14, Column Name: Steel_Plate_Thickness
Index: 15, Column Name: Edges_Index
Index: 16, Column Name: Empty_Index
Index: 17, Column Name: Square_Index
Index: 18, Column Name: Outside_X_Index
Index: 19, Column Name: Edges_X_Index
Index: 20, Column Name: Edges_Y_Index
Index: 21, Column Name: Outside_Global_Index
Index: 22, Column Name: LogOfAreas
Index: 23, Column Name: Log_X_Index
Index: 24, Column Name: Log_Y_Index
Index: 25, Column Name: Orientation_Index
Index: 26, Column Name: Luminosity_Index
Index: 27, Column Name: SigmoidOfAreas
Index: 28, Column Name: Pastry
Index: 29, Column Name: Z_Scratch
Index: 30, Column Name: K_Scratch
Index: 31, Column Name: Stains
Index: 32, Column Name: Dirtiness
Index: 33, Column Name: Bumps
Index: 34, Column Name: Other_Faults

```

Dessas, as nossas *features* são as 27 primeiras e os *labels* são as restantes:

```

[9]: features = faults.columns[:27]
labels = faults.columns[27:]

X = faults[features]
Y = faults[labels]

```

1.0.2 2. Faça uma análise exploratória dos dados.

```
[10]: X.describe()
```

```
[10]: shape: (9, 28)
```

statistic	X_Minimum	X_Maximum	Y_Minimum	...	Log_Y_Ind	Orientati
Luminosit	Sigmoid0					
---	---	---	---		ex	on_Index
y_Index	fAreas					
str	f64	f64	f64		---	---
---	---					
					f64	f64
f64	f64					
count	1941.0	1941.0	1941.0	...	1941.0	1941.0

1941.0	1941.0					
null_coun	0.0	0.0	0.0	...	0.0	0.0
0.0	0.0					
t						
mean	571.13601	617.96445	1.6507e6	...	1.403271	0.083288
-0.131305	0.58542					
	2	1				
std	520.69067	497.62741	1.7746e6	...	0.454345	0.500868
0.148767	0.339452					
	1					
min	0.0	4.0	6712.0	...	0.0	-0.991
-0.9989	0.119					
25%	51.0	192.0	471253.0	...	1.0792	-0.3333
-0.195	0.2482					
50%	435.0	467.0	1.204128e	...	1.3222	0.0952
-0.133	0.5063					
			6			
75%	1053.0	1072.0	2.183073e	...	1.7324	0.5116
-0.0666	0.9998					
			6			
max	1705.0	1713.0	1.2987661	...	4.2587	0.9917
0.6421	1.0					
			e7			

i. A base de dados é consistente? Sim! Os tipos de cada coluna estão bem definidos e organizados.

ii. Há dados faltantes? Não. Podemos ver dos dataframes gerados acima (X e Y) que não existem dados faltantes (`null_count = 0` em todas as colunas).

iii. Há dados não numéricos? Não. As características das placas de aço (dataframe X) não contém dados não numéricos. Já os defeitos (dataframe Y), que são categorias, foram organizados no formato *one-hot encoding*. Além disso, há também duas colunas em X codificadas com *one-hot encoding*, `TypeOfSteel_A300`, `TypeOfSteel_A400`.

iv. A base de dados é balanceada?

```
[11]: faults_count = Y.sum()
```

```
faults_count
```

```
[11]: shape: (1, 7)
```

Pastry	Z_Scratch	K_Scratch	Stains	Dirtiness	Bumps	Other_Faults
---	---	---	---	---	---	---
i64	i64	i64	i64	i64	i64	i64
158	190	391	72	55	402	673


Como se vê acima, a base de dados não é balanceada quanto às classes (tipos de defeitos).


v. As variáveis assumem valores plausíveis? Como é a distribuição dos dados? Há outliers? Faça boxplots e comente. Segundo o output dos métodos describe acima, os dados parecem assumir valores plausíveis. Vamos checar isso com boxplots. Removendo as colunas *one-hot encoding* do dataframe X, temos:

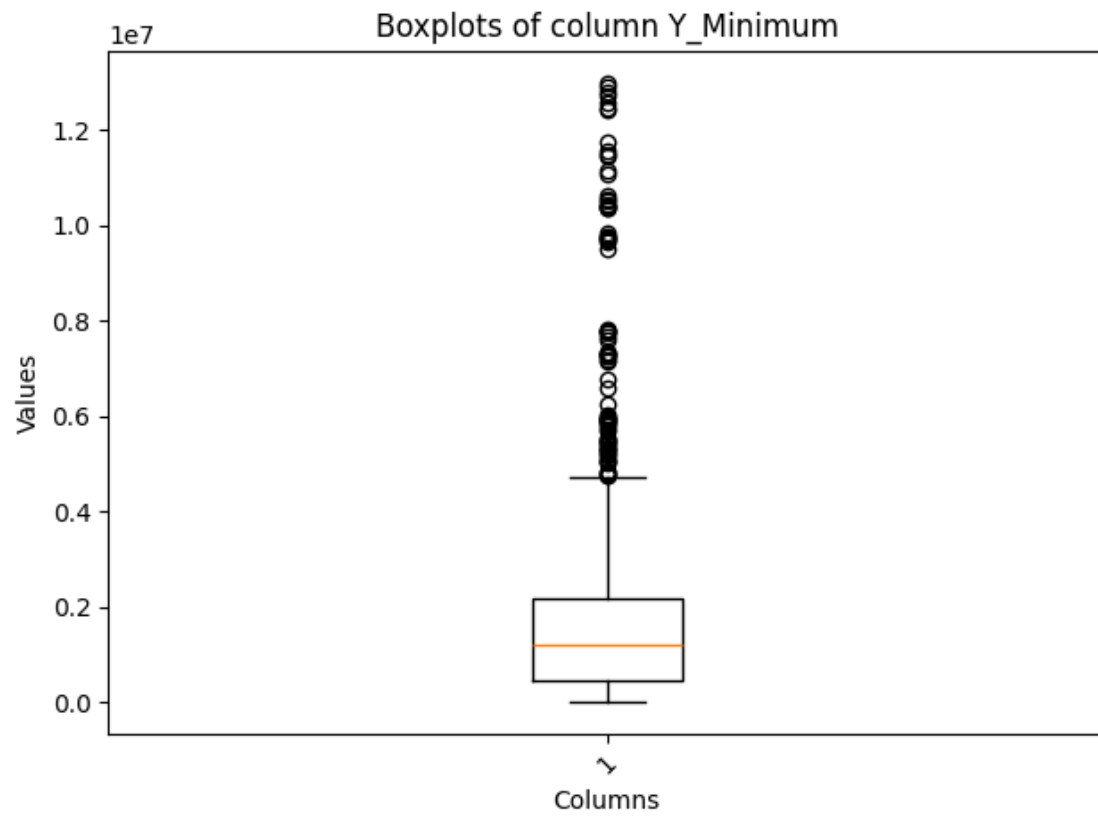
```
[12]: X_without_steel_type = X.select(pl.selectors.exclude('TypeOfSteel_A300',  
    ↪ 'TypeOfSteel_A400'))  
X_for_plotting = np.array([X_without_steel_type[col].to_numpy() for col in  
    ↪ X_without_steel_type.columns])  
  
X_for_plotting.shape
```

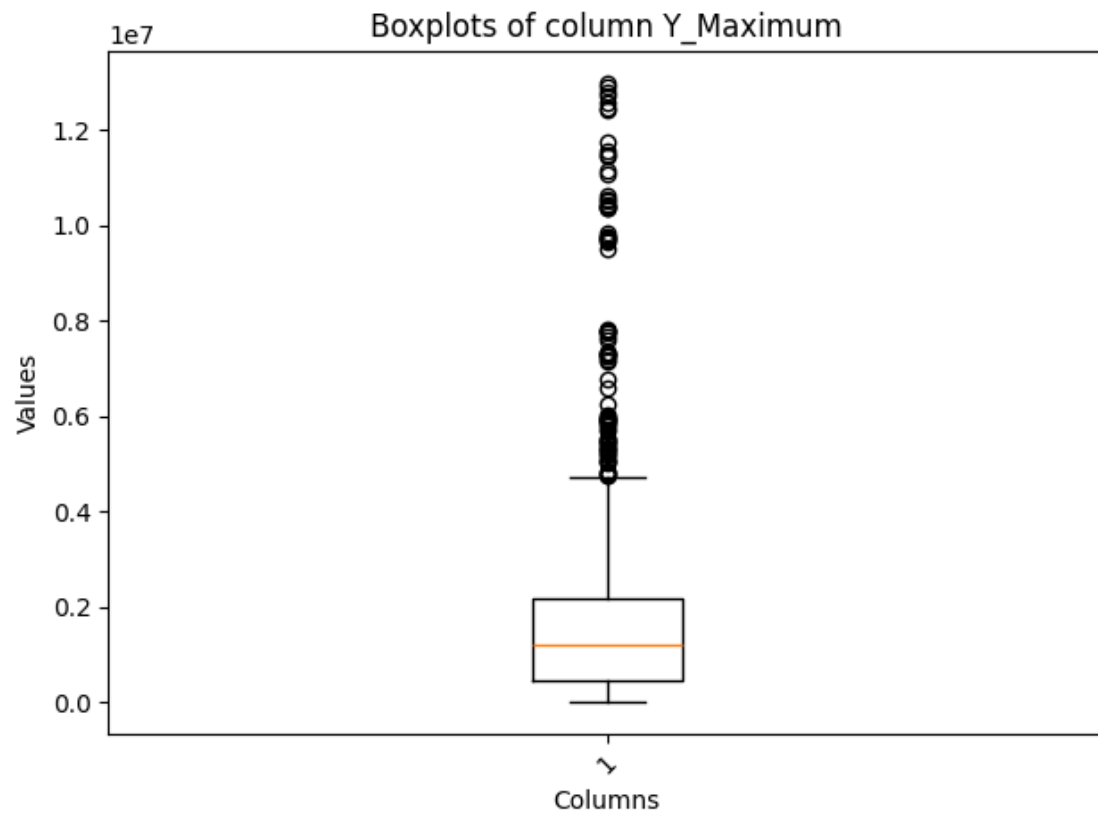
```
[12]: (25, 1941)
```

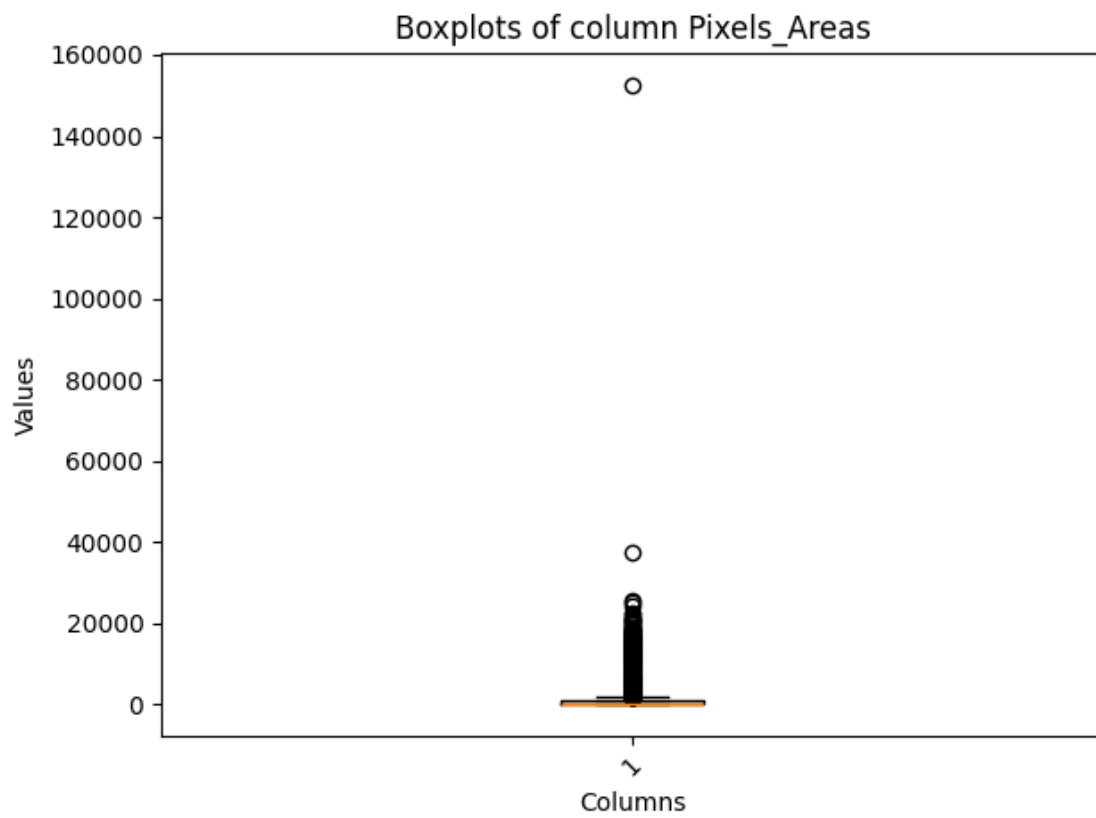
```
[13]: for i in range(X_without_steel_type.width):  
    plt.boxplot(X_for_plotting[i])  
    plt.title(f"Boxplots of column {X_without_steel_type.columns[i]}")  
    plt.ylabel("Values")  
    plt.xlabel("Columns")  
    plt.xticks(rotation=45)  
    plt.tight_layout()  
    plt.show();
```

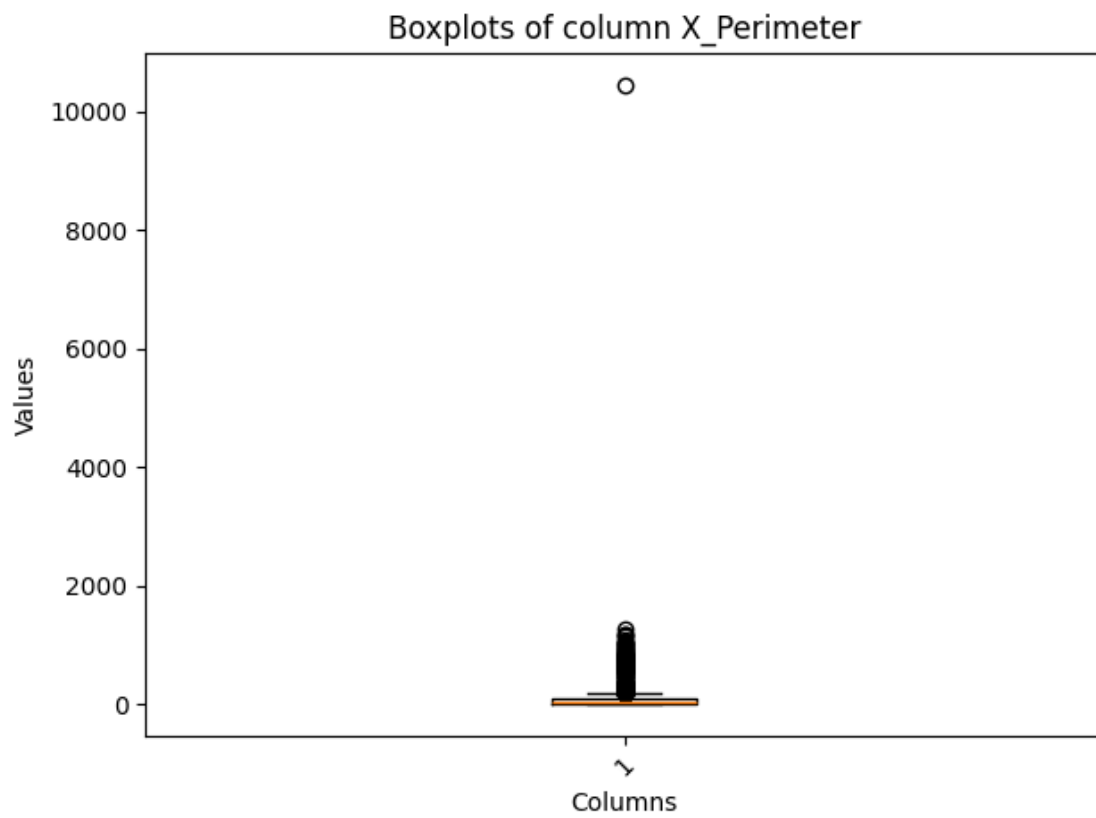


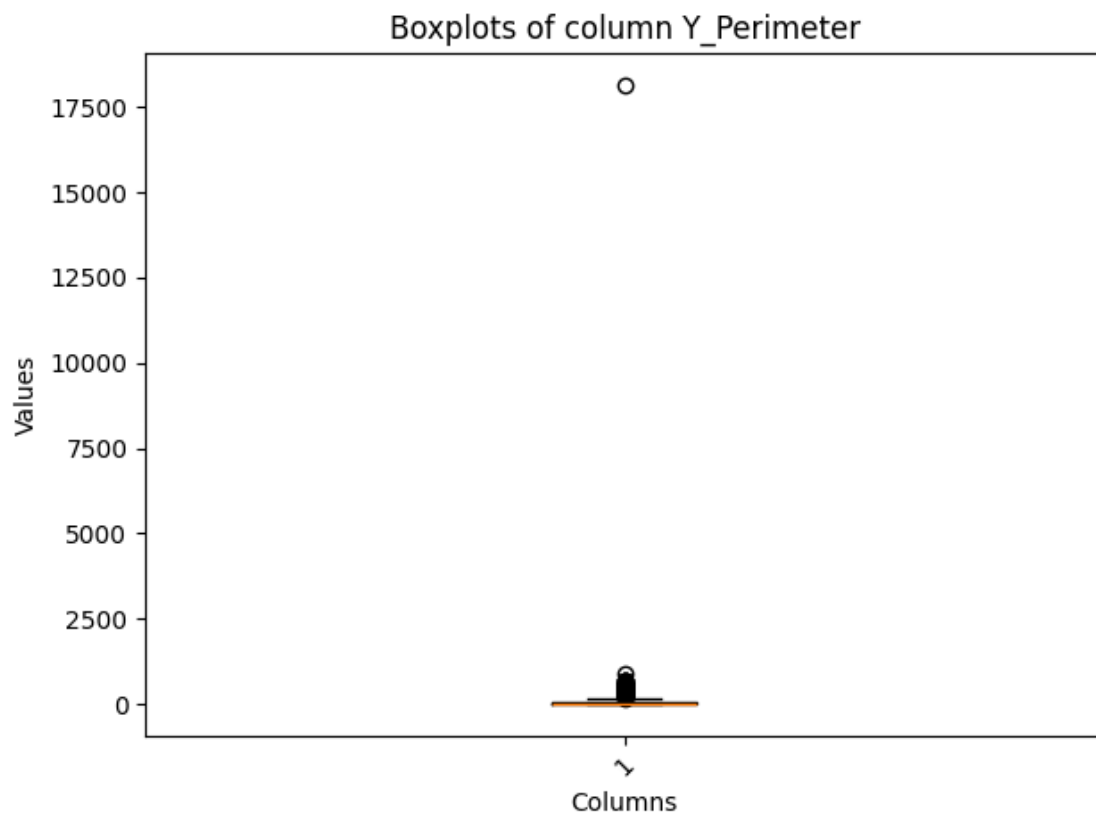




















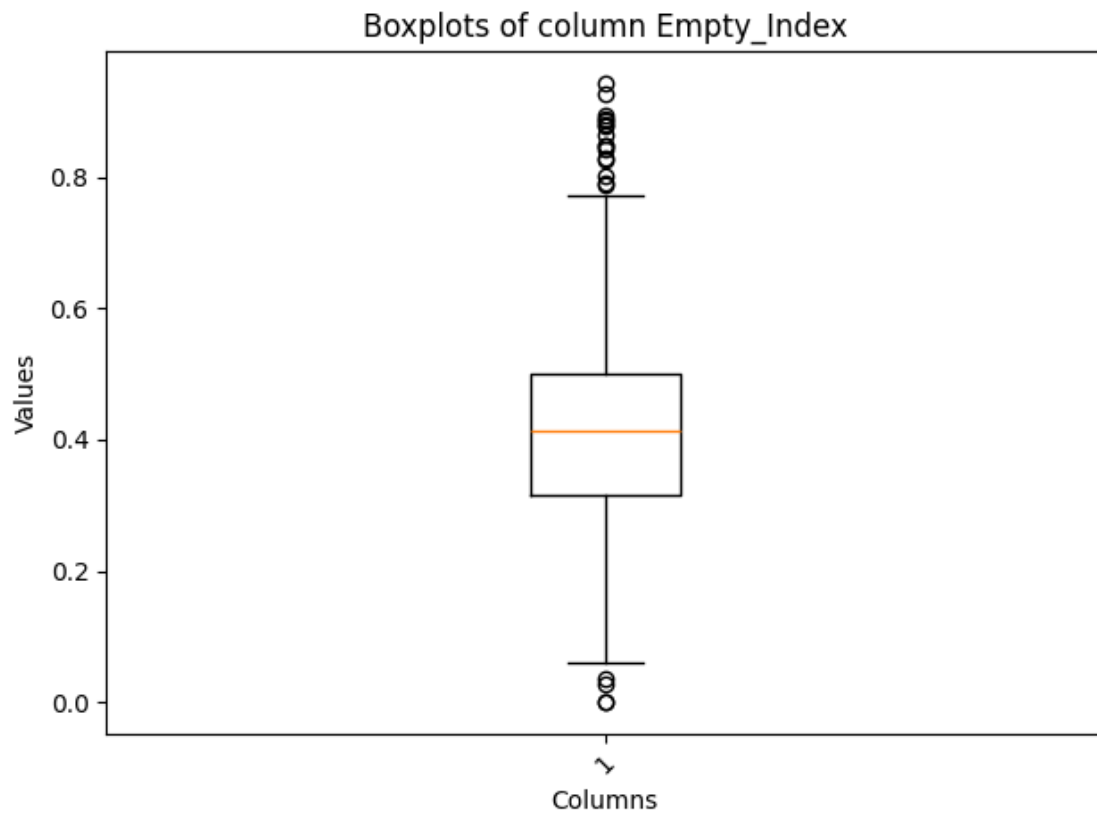





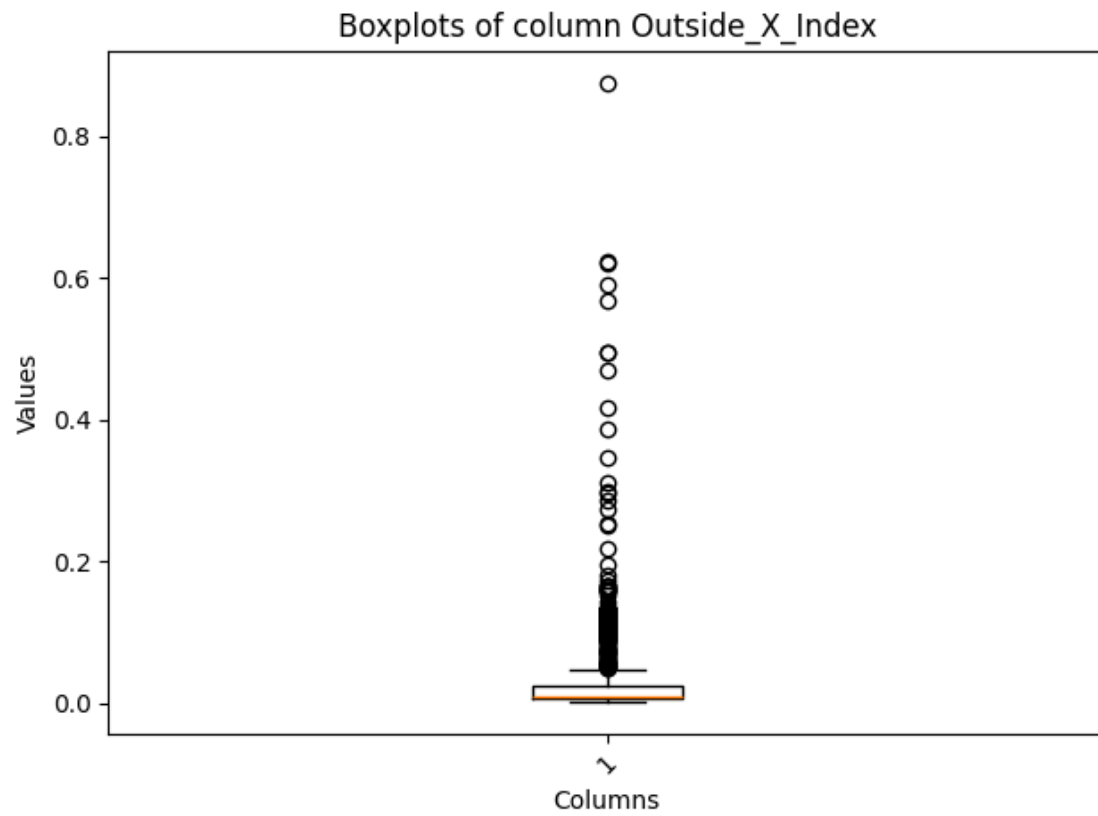


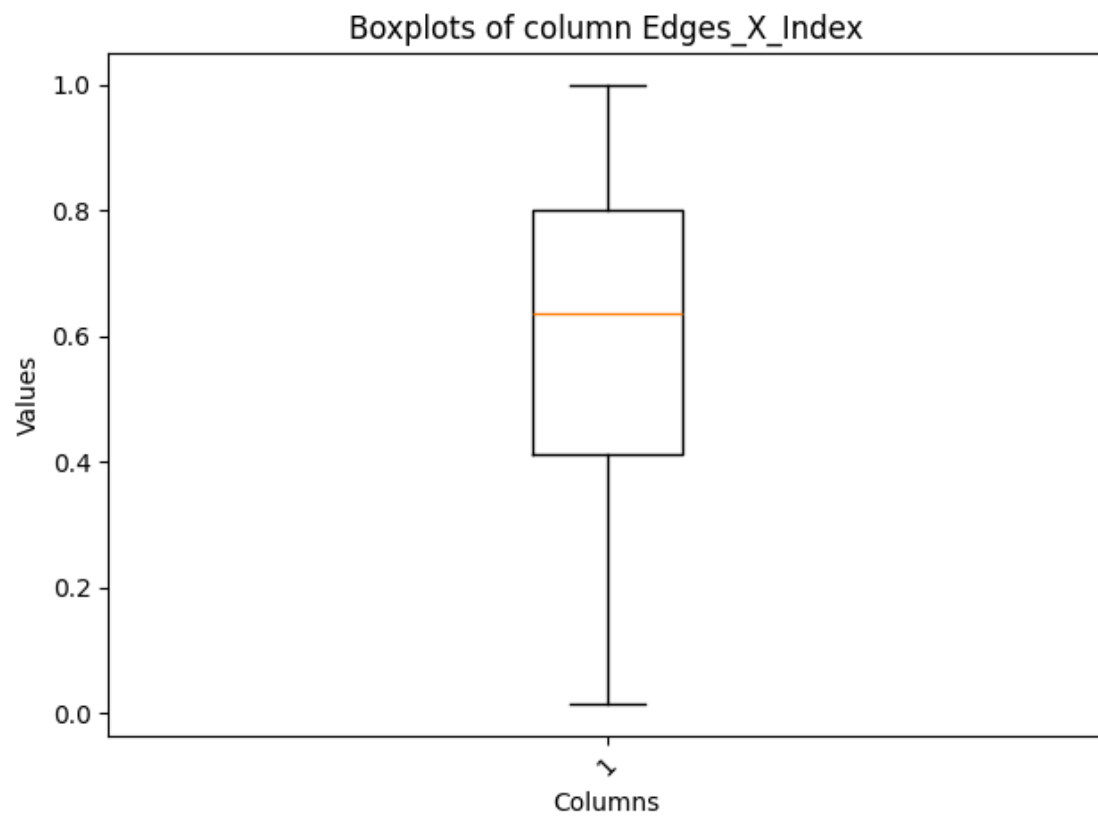


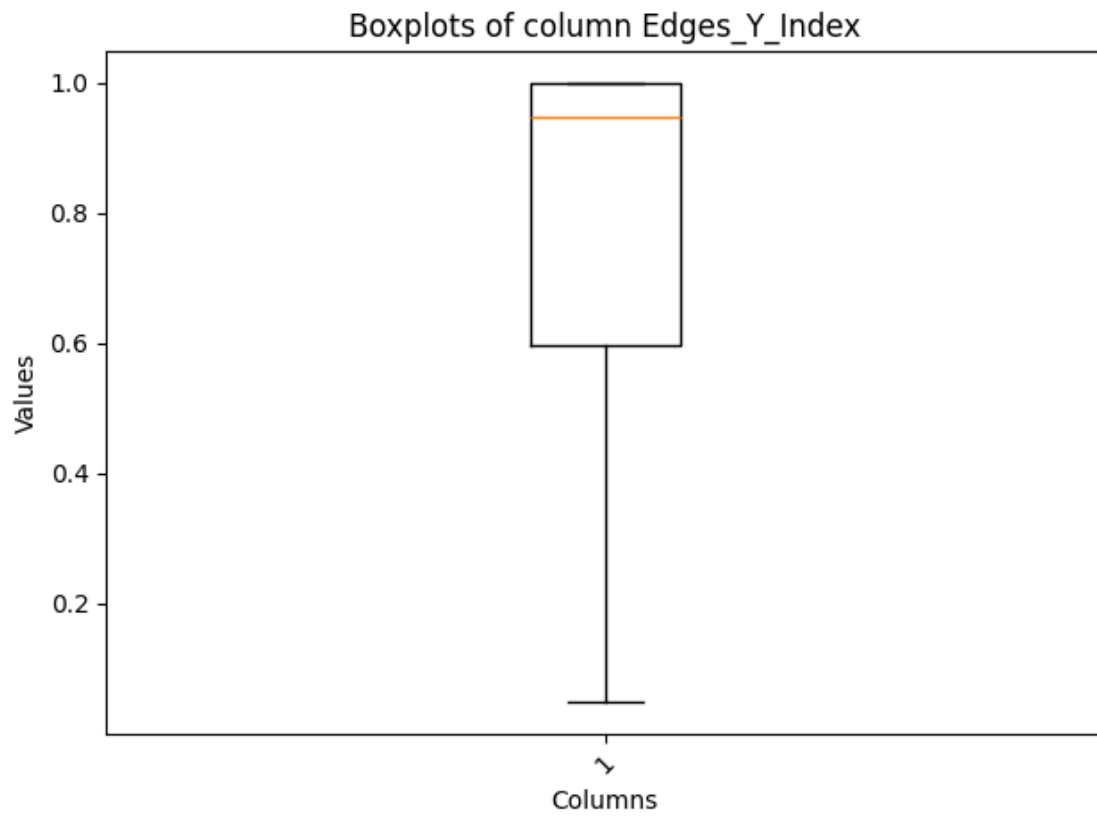


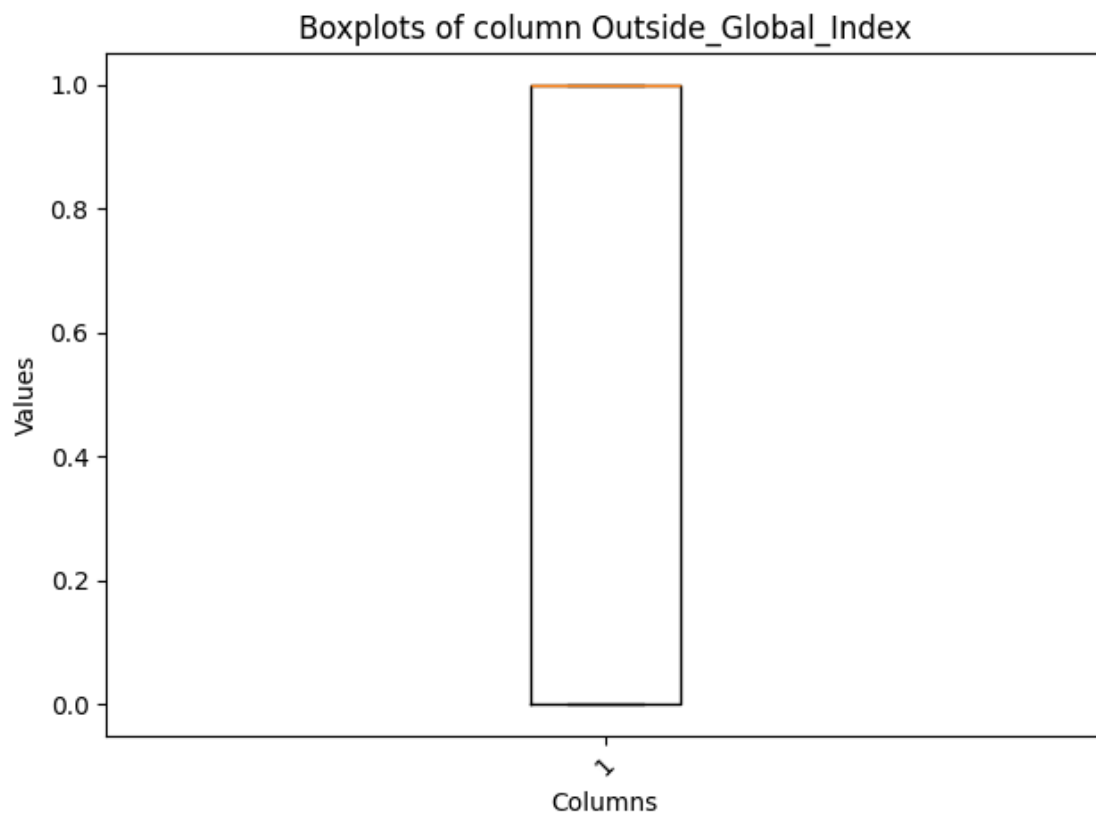











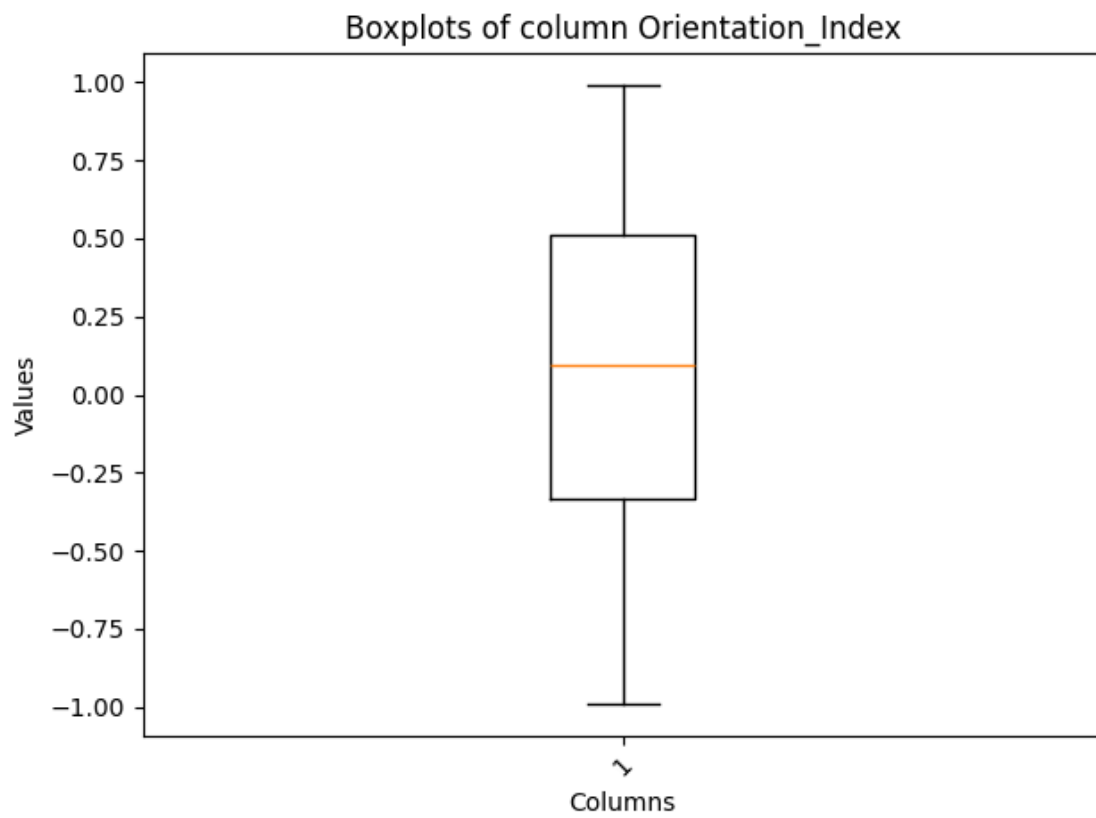





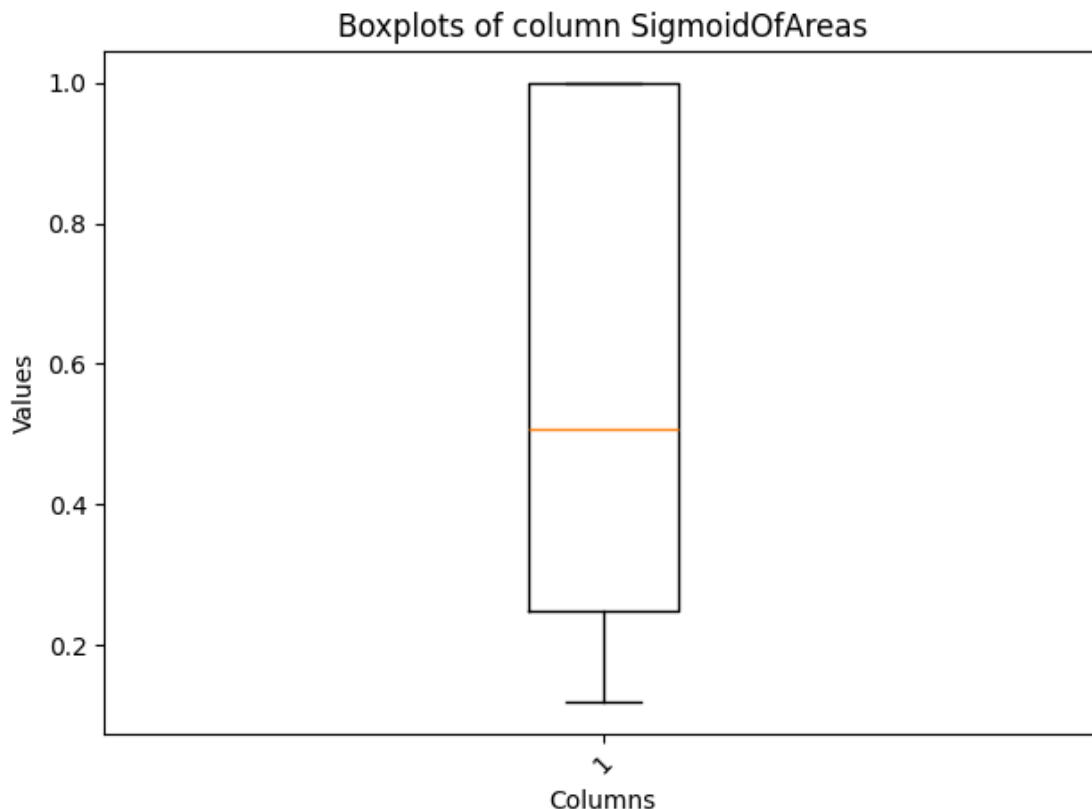












Parece que existem muitos outliers em várias colunas, vamos quantificá-los:

```
[14]: def calculate_outliers_percentage(series):
    q1 = series.quantile(0.25)
    q3 = series.quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - 1.5 * iqr
    upper_bound = q3 + 1.5 * iqr

    # Count outliers
    outliers_count = series.filter((series < lower_bound) | (series >
    ↪upper_bound)).shape[0]
    total_count = series.shape[0]

    # Calculate percentage of outliers
    percentage = (outliers_count / total_count) * 100
    return round(percentage, 2)

[15]: outlier_percentages = {col:
    ↪calculate_outliers_percentage(X_without_steel_type[col]) for col in
    ↪X_without_steel_type.columns}
```



```

outlier_percentages = {'statistic': 'outlier percentage'} | outlier_percentages

# Create a new DataFrame for outlier percentages
outlier_row = pl.DataFrame(outlier_percentages)

# Add the outlier percentages row to the original DataFrame
X_description_with_outliers = pl.concat([X_without_steel_type.describe(),
    outlier_row])

X_description_with_outliers

```

[15]: shape: (10, 26)

statistic	X_Minimum	X_Maximum	Y_Minimum	...	Log_Y_Ind	Orientati
Luminosit	Sigmoid0					
---	---	---	---		ex	on_Index
y_Index	fAreas					
str	f64	f64	f64		---	---
---	---					
					f64	f64
f64	f64					
count	1941.0	1941.0	1941.0	...	1941.0	1941.0
1941.0	1941.0					
null_coun	0.0	0.0	0.0	...	0.0	0.0
0.0	0.0					
t						
mean	571.13601	617.96445	1.6507e6	...	1.403271	0.083288
-0.131305	0.58542					
	2	1				
std	520.69067	497.62741	1.7746e6	...	0.454345	0.500868
0.148767	0.339452					
	1					
min	0.0	4.0	6712.0	...	0.0	-0.991
-0.9989	0.119					
25%	51.0	192.0	471253.0	...	1.0792	-0.3333
-0.195	0.2482					
50%	435.0	467.0	1.204128e	...	1.3222	0.0952
-0.133	0.5063					
			6			
75%	1053.0	1072.0	2.183073e	...	1.7324	0.5116

```

-0.0666    0.9998
6
max      1705.0    1713.0    1.2987661 ... 4.2587    0.9917
0.6421    1.0
e7
outlier    0.0      0.0      4.17      ... 0.21      0.0
6.9      0.0
percentag
e

```

```

[16]: outliers_row = X_description_with_outliers.row(by_predicate = (pl.
    ↪col('statistic') == "outlier percentage"))

outlier_percentage_threshold = 15
big_outliers_indexes = [
    index for index, value in enumerate(outliers_row)
    if isinstance(value, (int, float)) and value > outlier_percentage_threshold
]
X_description_big_outliers = X_description_with_outliers[:,
    ↪big_outliers_indexes]

X_description_big_outliers

```

```

[16]: shape: (10, 4)

```

Pixels_Areas	X_Perimeter	Sum_of_Luminosity	Outside_X_Index
---	---	---	---
f64	f64	f64	f64
1941.0	1941.0	1941.0	1941.0
0.0	0.0	0.0	0.0
1893.878413	111.855229	206312.147862	0.033361
5168.45956	301.209187	512293.587609	0.058961
2.0	2.0	250.0	0.0015
84.0	15.0	9522.0	0.0066
174.0	26.0	19202.0	0.0101
822.0	84.0	83011.0	0.0235
152655.0	10449.0	1.1591414e7	0.8759
20.35	18.13	20.56	19.06

```
[17]: X_description_big_outliers.columns
```

```
[17]: ['Pixels_Areas', 'X_Perimeter', 'Sum_of_Luminosity', 'Outside_X_Index']
```

Como se vê, as características (*features*) acima possuem um alto número de outliers (mais do que 15%).

vi. Há necessidade de normalizar ou padronizar as variáveis de entrada? Justifique. Parece que sim, posto que há uma variabilidade muito grande entre os valores de cada variável (indo da ordem de 10 até 10) e sabemos que redes neurais funcionam melhor com dados mais uniformes. Vamos fazer isso posteriormente, antes do treinamento.

vii. Analise o heatmap das variáveis e proponha uma estratégia para reduzir as variáveis de entrada sem perda de informação útil para o classificador. Heatmap com o matplotlib, excluindo as variáveis (*features*) codificadas com *one-hot encoding*:

```
[18]: X_without_steel_type.corr()
```

```
[18]: shape: (25, 25)
```


X_Minimum	X_Maximum	Y_Minimum	Y_Maximum	...	Log_Y_Ind	Orientati
Luminosit	Sigmoid0					
---	---	---	---		ex	on_Index
y_Index	fAreas					
f64	f64	f64	f64		---	---
---	---					
					f64	f64
f64	f64					
1.0	0.988314	0.041821	0.041807	...	-0.326851	0.178585
-0.031578	-0.35525					
1						
0.988314	1.0	0.052147	0.052135	...	-0.26599	0.115019
-0.038996	-0.28673					
6						
0.041821	0.052147	1.0	1.0	...	-0.008442	-0.086497
-0.090654	0.025257					
0.041807	0.052135	1.0	1.0	...	-0.008382	-0.08648
-0.090666	0.025284					
-0.307322	-0.225399	0.01767	0.01784	...	0.578342	-0.137604
-0.043449	0.422947					
...
...						

-0.437944	-0.324012	0.070406	0.070432	...	0.598652	-0.536629
-0.064923	0.757343					
-0.326851	-0.26599	-0.008442	-0.008382	...	1.0	0.316792
-0.21911	0.838188					
0.178585	0.115019	-0.086497	-0.08648	...	0.316792	1.0
-0.153464	-0.02397					
8						
-0.031578	-0.038996	-0.090654	-0.090666	...	-0.21911	-0.153464
1.0	-0.18484					
-0.355251	-0.286736	0.025257	0.025284	...	0.838188	-0.023978
-0.18484	1.0					

```
[19]: import seaborn as sns

dim = (11.7, 8.27)
fig, ax = plt.subplots(figsize = dim)

ax = sns.heatmap(
    X_without_steel_type.corr(),
    linewidth=0.5,
    cmap='crest',
    xticklabels=X_without_steel_type.columns, yticklabels=X_without_steel_type.
    ↪columns)
plt.show()
```



Visualmente podemos perceber que existe correlação mais forte entre as variáveis X_Minimum/X_Maximum, Y_Minimum/Y_Maximum, Pixels_Areas/X_Perimeter/Y_Perimeter/Sum_of_Luminosity, LogOfAreas/Log_X_Index/Log_Y_Index. Vamos quantificar melhor a magnitude dessas correlações, considerando apenas variáveis cuja correlação é maior do que 95%:

```
[20]: correlation_threshold = .95
X_corr = X_without_steel_type.corr()
corr_var_indexes = [
    {i, j} for i in range(X_corr.height) for j in range(X_corr.width)
    if abs(X_corr[i, j]) > correlation_threshold
    and i != j
]
corr_var_indexes = list({frozenset(i) for i in corr_var_indexes})

# Display the results
for row, col in corr_var_indexes:
    print(f"Value at ({X_corr.columns[row]}, '{X_corr.columns[col]}') fulfills_
    ↳the condition.")
```

Value at (Pixels_Areas, 'X_Perimeter') fulfills the condition.
 Value at (X_Minimum, 'X_Maximum') fulfills the condition.
 Value at (Y_Minimum, 'Y_Maximum') fulfills the condition.
 Value at (Pixels_Areas, 'Sum_of_Luminosity') fulfills the condition.

Como se vê acima, os pares de variáveis listados têm uma correlação de mais de 95%. Vamos dropar as variáveis redundantes:

```
[21]: X = X.drop(['X_Perimeter', 'X_Maximum', 'Y_Maximum', 'Sum_of_Luminosity'],
↳strict = False)
```

X

```
[21]: shape: (1_941, 23)
```

X_Minimum	Y_Minimum	Pixels_Ar	Y_Perimet	...	Log_Y_Ind	Orientati
Luminosit	Sigmoid0					
---	---	eas	er		ex	on_Index
y_Index	fAreas					
i64	i64	---	---		---	---
---	---					
		i64	i64		f64	f64
f64	f64					
42	270900	267	44	...	1.6435	0.8182
-0.2913	0.5822					
645	2538079	108	30	...	1.4624	0.7931
-0.1756	0.2984					
829	1553913	71	19	...	1.2553	0.6667
-0.1228	0.215					
853	369370	176	45	...	1.6532	0.8444
-0.1568	0.5212					
1289	498078	2409	260	...	2.4099	0.9338
-0.1992	1.0					
...
...						
249	325780	273	22	...	1.2041	-0.4286
0.0026	0.7254					
144	340581	287	24	...	1.2305	-0.4516
-0.0582	0.8173					
145	386779	292	22	...	1.1761	-0.4828
0.0052	0.7079					
137	422497	419	47	...	1.4914	-0.0606
-0.0171	0.9919					
1261	87951	103	22	...	1.2041	-0.2
-0.1139	0.5296					

Vamos salvar nossos dados nas variáveis `X_0` e `Y_0` antes de manipulá-los mais:

```
[22]: X_0 = X
      Y_0 = Y
```

1.0.3 3. Particione aleatoriamente 70% das amostras para treinar a rede neural MLP e o restante das amostras para validar o sistema.

Antes de proceder com a partição do conjunto, vamos aplicar algumas transformações nos dados visando melhorar o resultado da rede neural.

```
[47]: # Retrieving the data from the saved state in the code cell above

X = X_0
Y = Y_0
```

Repare no entanto que nos dados de entrada existem *features* categóricas que *podem* não ser transformadas corretamente. Vamos salvá-las a parte:

```
[48]: # Lets remove the categorical features before winsorization
features_categorical = ["TypeOfSteel_A300", "TypeOfSteel_A400"]
X_categorical = X[features_categorical]

X_continuous = X.select(pl.col("*").exclude(features_categorical))
features_continuous = X_continuous.columns
```

Vamos primeiramente remover os valores mais discrepantes através da função `winsorize` ([link da documentação](#)), que limita os valores limites nos dados, reduzindo o efeito dos *outliers* (ver *Winsorizing*):

```
[49]: from scipy.stats import mstats

# Winsorize the data (e.g., cap at 5st and 95th percentiles)
X_winsorized = mstats.winsorize(X_continuous.to_numpy(), limits=[0.05, 0.05],
                                ↪axis = 0)
X_winsorized = pl.DataFrame(X_winsorized, schema = features_continuous)
X_continuous = X_winsorized
```

Podemos perceber que há uma diminuição substantiva na variabilidade dos dados após o processo acima. Vamos agora normalizar as variáveis de entrada com auxílio do pacote `sklearn` ([link](#)):

```
[50]: from sklearn.preprocessing import StandardScaler, RobustScaler, MinMaxScaler

# Initialize the scaler
scaler = MinMaxScaler(feature_range = (-1, 1))
```

```
# Fit and transform the features
X_scaled = scaler.fit_transform(X_continuous)
X_continuous = pl.DataFrame(X_scaled, schema = features_continuous)
```

Antes de prosseguir com a partição dos dados, vamos tratar do problema de balanceamento das classes. Conforme visto anteriormente, a distribuição é:

```
[51]: Y.sum()
```

```
[51]: shape: (1, 7)
```

Pastry	Z_Scratch	K_Scratch	Stains	Dirtiness	Bumps	Other_Faults
---	---	---	---	---	---	---
i64	i64	i64	i64	i64	i64	i64
158	190	391	72	55	402	673

Existem duas possíveis abordagens para resolver isso: *oversampling* e *undersampling*. De maneira ingênua, podemos assumir que a última técnica desprezaria dados, nos fazendo perder informação. Vamos então adotar o *oversampling* com auxílio da biblioteca [imbalanced-learn](#).

```
[52]: X = pl.concat([X_continuous, X_categorical], how = "horizontal")
```

Vamos salvar nossos dados antes do balanceamento para posteriormente demonstrarmos a importância dessa etapa:

```
[53]: X_unbalanced = X
Y_unbalanced = Y
```

```
[54]: from imblearn.over_sampling import SMOTE, ADASYN, SMOTENC

# smote_nc = SMOTENC(categorical_features = [21, 22])
# X_over, Y_over = smote_nc.fit_resample(X.to_numpy(), Y.to_numpy())

X_over, Y_over = SMOTE().fit_resample(X.to_numpy(), Y.to_numpy())
```

```
[55]: Y = pl.DataFrame(Y_over, schema = Y.columns)
X = pl.DataFrame(X_over, schema = X.columns)

Y.sum()
```

```
[55]: shape: (1, 7)
```

Pastry	Z_Scratch	K_Scratch	Stains	Dirtiness	Bumps	Other_Faults
---	---	---	---	---	---	---
i64	i64	i64	i64	i64	i64	i64

673

673

673

673

673

673

673

Particionando em 70% a proporção para treino e teste. Repare no argumento `shuffle` como `True` que garante que haverá embaralhamento dos dados, que inicialmente estavam organizados por falha (primeiro todas as falhas do tipo *Pastry*, depois do tipo *Z_Scratch*, etc):

```
[56]: from sklearn.model_selection import train_test_split

train_fraction = 0.7
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = (1 -
↪train_fraction), shuffle = True)
```

1.0.4 4. Implemente uma rede neural MLP capaz de receber as variáveis de entrada que descrevem a geometria do defeito na placa e indicar a probabilidade de existência de cada possível falha.

Um modelo simples de MLP que é capaz de receber as variáveis e gerar o rótulo correspondente é:

```
[57]: from tensorflow import keras
from keras.models import Sequential
from keras.layers import Flatten
from keras.layers import Dense
from keras.optimizers import Adam

keras.backend.clear_session()

model = Sequential()
model.add(Dense(200, activation='relu')),
model.add(Dense(50, activation='relu')),
model.add(Dense(10, activation='relu')),
model.add(Dense(Y_train.width, activation='softmax'))
```

```
[58]: optimizer = Adam()
model.compile(optimizer=optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
[59]: history = model.fit(
    X_train.to_numpy(),
    Y_train.to_numpy(),
    epochs = 128,
    batch_size = 32,
    validation_data = (X_test.to_numpy(), Y_test.to_numpy()))
```

Epoch 1/128

104/104

1s 3ms/step -

accuracy: 0.4824 - loss: 1.5507 - val_accuracy: 0.7256 - val_loss: 0.8164
 Epoch 2/128
 104/104 0s 1ms/step -
 accuracy: 0.7736 - loss: 0.7199 - val_accuracy: 0.8154 - val_loss: 0.5819
 Epoch 3/128
 104/104 0s 2ms/step -
 accuracy: 0.8176 - loss: 0.5444 - val_accuracy: 0.8161 - val_loss: 0.5626
 Epoch 4/128
 104/104 0s 1ms/step -
 accuracy: 0.8255 - loss: 0.5044 - val_accuracy: 0.8303 - val_loss: 0.4821
 Epoch 5/128
 104/104 0s 1ms/step -
 accuracy: 0.8535 - loss: 0.4042 - val_accuracy: 0.8423 - val_loss: 0.4407
 Epoch 6/128
 104/104 0s 1ms/step -
 accuracy: 0.8625 - loss: 0.3679 - val_accuracy: 0.8501 - val_loss: 0.4307
 Epoch 7/128
 104/104 0s 1ms/step -
 accuracy: 0.8758 - loss: 0.3555 - val_accuracy: 0.8444 - val_loss: 0.4208
 Epoch 8/128
 104/104 0s 1ms/step -
 accuracy: 0.8774 - loss: 0.3267 - val_accuracy: 0.8557 - val_loss: 0.4004
 Epoch 9/128
 104/104 0s 1ms/step -
 accuracy: 0.8696 - loss: 0.3573 - val_accuracy: 0.8586 - val_loss: 0.4058
 Epoch 10/128
 104/104 0s 1ms/step -
 accuracy: 0.8897 - loss: 0.3108 - val_accuracy: 0.8628 - val_loss: 0.3721
 Epoch 11/128
 104/104 0s 2ms/step -
 accuracy: 0.8960 - loss: 0.2870 - val_accuracy: 0.8593 - val_loss: 0.3834
 Epoch 12/128
 104/104 0s 2ms/step -
 accuracy: 0.8996 - loss: 0.2914 - val_accuracy: 0.8621 - val_loss: 0.3731
 Epoch 13/128
 104/104 0s 2ms/step -
 accuracy: 0.8975 - loss: 0.2782 - val_accuracy: 0.8755 - val_loss: 0.3520
 Epoch 14/128
 104/104 0s 1ms/step -
 accuracy: 0.8995 - loss: 0.2652 - val_accuracy: 0.8762 - val_loss: 0.3556
 Epoch 15/128
 104/104 0s 2ms/step -
 accuracy: 0.9094 - loss: 0.2674 - val_accuracy: 0.8649 - val_loss: 0.3900
 Epoch 16/128
 104/104 0s 2ms/step -
 accuracy: 0.8976 - loss: 0.2734 - val_accuracy: 0.8564 - val_loss: 0.3916
 Epoch 17/128
 104/104 0s 1ms/step -

accuracy: 0.8718 - loss: 0.3442 - val_accuracy: 0.8777 - val_loss: 0.3496
 Epoch 18/128
 104/104 0s 2ms/step -
 accuracy: 0.9168 - loss: 0.2285 - val_accuracy: 0.8713 - val_loss: 0.3903
 Epoch 19/128
 104/104 0s 1ms/step -
 accuracy: 0.8950 - loss: 0.3064 - val_accuracy: 0.8833 - val_loss: 0.3404
 Epoch 20/128
 104/104 0s 2ms/step -
 accuracy: 0.9199 - loss: 0.2275 - val_accuracy: 0.8847 - val_loss: 0.3503
 Epoch 21/128
 104/104 0s 1ms/step -
 accuracy: 0.9162 - loss: 0.2169 - val_accuracy: 0.8840 - val_loss: 0.3389
 Epoch 22/128
 104/104 0s 2ms/step -
 accuracy: 0.9324 - loss: 0.2027 - val_accuracy: 0.8911 - val_loss: 0.3320
 Epoch 23/128
 104/104 0s 2ms/step -
 accuracy: 0.9178 - loss: 0.2221 - val_accuracy: 0.8911 - val_loss: 0.3261
 Epoch 24/128
 104/104 0s 2ms/step -
 accuracy: 0.9319 - loss: 0.1912 - val_accuracy: 0.8897 - val_loss: 0.3355
 Epoch 25/128
 104/104 0s 2ms/step -
 accuracy: 0.9281 - loss: 0.1924 - val_accuracy: 0.8868 - val_loss: 0.3421
 Epoch 26/128
 104/104 0s 2ms/step -
 accuracy: 0.9346 - loss: 0.1748 - val_accuracy: 0.8946 - val_loss: 0.3396
 Epoch 27/128
 104/104 0s 2ms/step -
 accuracy: 0.9410 - loss: 0.1691 - val_accuracy: 0.8847 - val_loss: 0.3684
 Epoch 28/128
 104/104 0s 1ms/step -
 accuracy: 0.9278 - loss: 0.1908 - val_accuracy: 0.8996 - val_loss: 0.3211
 Epoch 29/128
 104/104 0s 1ms/step -
 accuracy: 0.9460 - loss: 0.1569 - val_accuracy: 0.8953 - val_loss: 0.3332
 Epoch 30/128
 104/104 0s 2ms/step -
 accuracy: 0.9509 - loss: 0.1549 - val_accuracy: 0.8748 - val_loss: 0.3699
 Epoch 31/128
 104/104 0s 2ms/step -
 accuracy: 0.9339 - loss: 0.1757 - val_accuracy: 0.9038 - val_loss: 0.3262
 Epoch 32/128
 104/104 0s 1ms/step -
 accuracy: 0.9510 - loss: 0.1524 - val_accuracy: 0.8812 - val_loss: 0.3463
 Epoch 33/128
 104/104 0s 2ms/step -

accuracy: 0.9428 - loss: 0.1544 - val_accuracy: 0.8982 - val_loss: 0.3402
 Epoch 34/128
 104/104 0s 2ms/step -
 accuracy: 0.9600 - loss: 0.1423 - val_accuracy: 0.8925 - val_loss: 0.3396
 Epoch 35/128
 104/104 0s 2ms/step -
 accuracy: 0.9577 - loss: 0.1305 - val_accuracy: 0.8967 - val_loss: 0.3335
 Epoch 36/128
 104/104 0s 2ms/step -
 accuracy: 0.9522 - loss: 0.1363 - val_accuracy: 0.8960 - val_loss: 0.3321
 Epoch 37/128
 104/104 0s 2ms/step -
 accuracy: 0.9482 - loss: 0.1380 - val_accuracy: 0.8939 - val_loss: 0.3731
 Epoch 38/128
 104/104 0s 2ms/step -
 accuracy: 0.9492 - loss: 0.1426 - val_accuracy: 0.8918 - val_loss: 0.3705
 Epoch 39/128
 104/104 0s 2ms/step -
 accuracy: 0.9431 - loss: 0.1662 - val_accuracy: 0.8678 - val_loss: 0.4032
 Epoch 40/128
 104/104 0s 2ms/step -
 accuracy: 0.9160 - loss: 0.2439 - val_accuracy: 0.8960 - val_loss: 0.3472
 Epoch 41/128
 104/104 0s 2ms/step -
 accuracy: 0.9648 - loss: 0.1113 - val_accuracy: 0.8967 - val_loss: 0.3483
 Epoch 42/128
 104/104 0s 2ms/step -
 accuracy: 0.9574 - loss: 0.1214 - val_accuracy: 0.9010 - val_loss: 0.3517
 Epoch 43/128
 104/104 0s 1ms/step -
 accuracy: 0.9626 - loss: 0.1178 - val_accuracy: 0.9010 - val_loss: 0.3315
 Epoch 44/128
 104/104 0s 2ms/step -
 accuracy: 0.9748 - loss: 0.0951 - val_accuracy: 0.8996 - val_loss: 0.3419
 Epoch 45/128
 104/104 0s 2ms/step -
 accuracy: 0.9610 - loss: 0.1117 - val_accuracy: 0.9045 - val_loss: 0.3437
 Epoch 46/128
 104/104 0s 2ms/step -
 accuracy: 0.9731 - loss: 0.0963 - val_accuracy: 0.8918 - val_loss: 0.3571
 Epoch 47/128
 104/104 0s 2ms/step -
 accuracy: 0.9713 - loss: 0.0937 - val_accuracy: 0.8996 - val_loss: 0.3496
 Epoch 48/128
 104/104 0s 2ms/step -
 accuracy: 0.9730 - loss: 0.0858 - val_accuracy: 0.9010 - val_loss: 0.3824
 Epoch 49/128
 104/104 0s 2ms/step -

accuracy: 0.9715 - loss: 0.0907 - val_accuracy: 0.8890 - val_loss: 0.3994
 Epoch 50/128
 104/104 0s 2ms/step -
 accuracy: 0.9578 - loss: 0.1196 - val_accuracy: 0.9017 - val_loss: 0.3429
 Epoch 51/128
 104/104 0s 2ms/step -
 accuracy: 0.9784 - loss: 0.0789 - val_accuracy: 0.8996 - val_loss: 0.3550
 Epoch 52/128
 104/104 0s 2ms/step -
 accuracy: 0.9764 - loss: 0.0784 - val_accuracy: 0.9088 - val_loss: 0.3567
 Epoch 53/128
 104/104 0s 2ms/step -
 accuracy: 0.9791 - loss: 0.0702 - val_accuracy: 0.9095 - val_loss: 0.3535
 Epoch 54/128
 104/104 0s 2ms/step -
 accuracy: 0.9790 - loss: 0.0754 - val_accuracy: 0.9081 - val_loss: 0.3625
 Epoch 55/128
 104/104 0s 2ms/step -
 accuracy: 0.9775 - loss: 0.0763 - val_accuracy: 0.9045 - val_loss: 0.3536
 Epoch 56/128
 104/104 0s 1ms/step -
 accuracy: 0.9749 - loss: 0.0728 - val_accuracy: 0.9024 - val_loss: 0.3780
 Epoch 57/128
 104/104 0s 1ms/step -
 accuracy: 0.9801 - loss: 0.0707 - val_accuracy: 0.8918 - val_loss: 0.3928
 Epoch 58/128
 104/104 0s 2ms/step -
 accuracy: 0.9406 - loss: 0.1726 - val_accuracy: 0.8996 - val_loss: 0.3969
 Epoch 59/128
 104/104 0s 2ms/step -
 accuracy: 0.9794 - loss: 0.0702 - val_accuracy: 0.9010 - val_loss: 0.4126
 Epoch 60/128
 104/104 0s 2ms/step -
 accuracy: 0.9756 - loss: 0.0793 - val_accuracy: 0.8989 - val_loss: 0.4238
 Epoch 61/128
 104/104 0s 2ms/step -
 accuracy: 0.9784 - loss: 0.0629 - val_accuracy: 0.9081 - val_loss: 0.3889
 Epoch 62/128
 104/104 0s 2ms/step -
 accuracy: 0.9797 - loss: 0.0627 - val_accuracy: 0.8960 - val_loss: 0.4228
 Epoch 63/128
 104/104 0s 1ms/step -
 accuracy: 0.9829 - loss: 0.0549 - val_accuracy: 0.8953 - val_loss: 0.4191
 Epoch 64/128
 104/104 0s 2ms/step -
 accuracy: 0.9851 - loss: 0.0541 - val_accuracy: 0.9059 - val_loss: 0.4194
 Epoch 65/128
 104/104 0s 1ms/step -

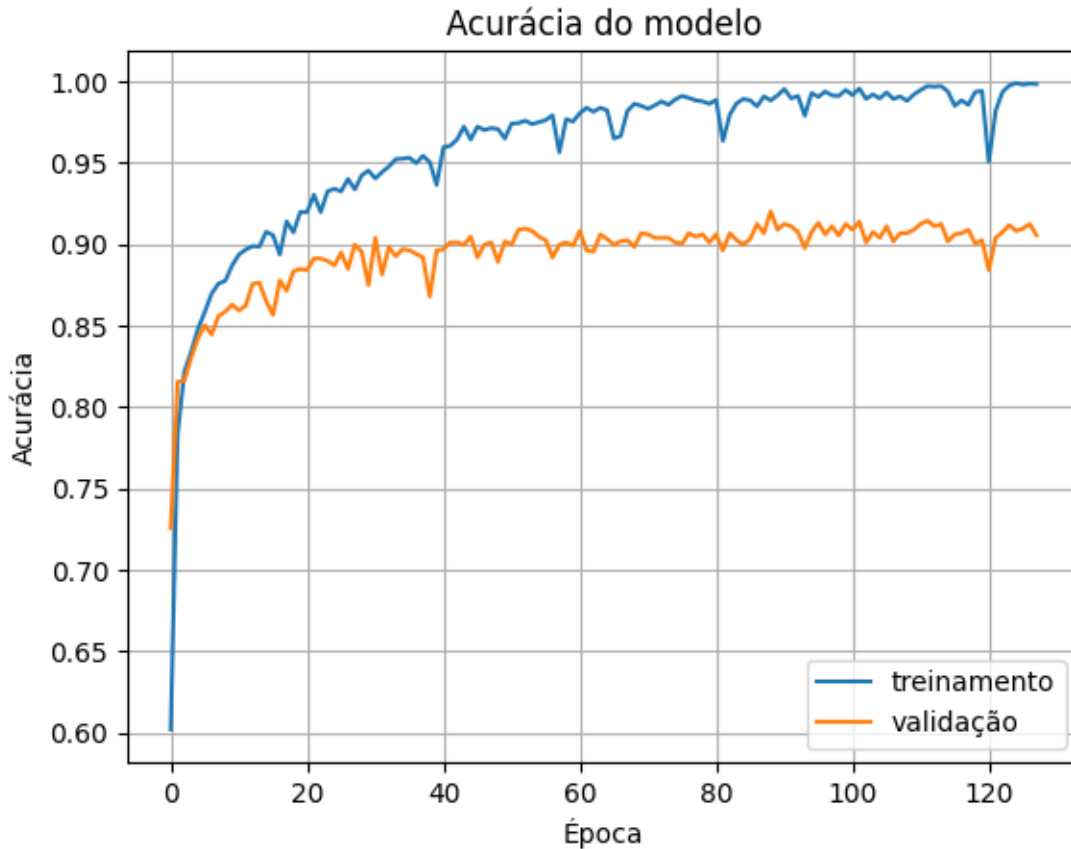
accuracy: 0.9865 - loss: 0.0461 - val_accuracy: 0.9031 - val_loss: 0.4011
 Epoch 66/128
 104/104 0s 1ms/step -
 accuracy: 0.9610 - loss: 0.1189 - val_accuracy: 0.8996 - val_loss: 0.4073
 Epoch 67/128
 104/104 0s 2ms/step -
 accuracy: 0.9528 - loss: 0.1543 - val_accuracy: 0.9017 - val_loss: 0.4106
 Epoch 68/128
 104/104 0s 1ms/step -
 accuracy: 0.9806 - loss: 0.0552 - val_accuracy: 0.9024 - val_loss: 0.4038
 Epoch 69/128
 104/104 0s 1ms/step -
 accuracy: 0.9871 - loss: 0.0467 - val_accuracy: 0.8982 - val_loss: 0.4386
 Epoch 70/128
 104/104 0s 2ms/step -
 accuracy: 0.9875 - loss: 0.0448 - val_accuracy: 0.9066 - val_loss: 0.4200
 Epoch 71/128
 104/104 0s 1ms/step -
 accuracy: 0.9814 - loss: 0.0532 - val_accuracy: 0.9059 - val_loss: 0.4079
 Epoch 72/128
 104/104 0s 2ms/step -
 accuracy: 0.9885 - loss: 0.0457 - val_accuracy: 0.9038 - val_loss: 0.4169
 Epoch 73/128
 104/104 0s 1ms/step -
 accuracy: 0.9881 - loss: 0.0437 - val_accuracy: 0.9038 - val_loss: 0.4307
 Epoch 74/128
 104/104 0s 2ms/step -
 accuracy: 0.9879 - loss: 0.0482 - val_accuracy: 0.9038 - val_loss: 0.4282
 Epoch 75/128
 104/104 0s 2ms/step -
 accuracy: 0.9900 - loss: 0.0474 - val_accuracy: 0.9010 - val_loss: 0.4333
 Epoch 76/128
 104/104 0s 2ms/step -
 accuracy: 0.9918 - loss: 0.0404 - val_accuracy: 0.9003 - val_loss: 0.4503
 Epoch 77/128
 104/104 0s 2ms/step -
 accuracy: 0.9905 - loss: 0.0388 - val_accuracy: 0.9066 - val_loss: 0.4555
 Epoch 78/128
 104/104 0s 2ms/step -
 accuracy: 0.9904 - loss: 0.0357 - val_accuracy: 0.9045 - val_loss: 0.4521
 Epoch 79/128
 104/104 0s 1ms/step -
 accuracy: 0.9873 - loss: 0.0470 - val_accuracy: 0.9059 - val_loss: 0.4474
 Epoch 80/128
 104/104 0s 1ms/step -
 accuracy: 0.9850 - loss: 0.0500 - val_accuracy: 0.9010 - val_loss: 0.4742
 Epoch 81/128
 104/104 0s 2ms/step -

accuracy: 0.9892 - loss: 0.0343 - val_accuracy: 0.9059 - val_loss: 0.4552
 Epoch 82/128
 104/104 0s 2ms/step -
 accuracy: 0.9575 - loss: 0.1261 - val_accuracy: 0.8960 - val_loss: 0.5016
 Epoch 83/128
 104/104 0s 2ms/step -
 accuracy: 0.9720 - loss: 0.0664 - val_accuracy: 0.9066 - val_loss: 0.4533
 Epoch 84/128
 104/104 0s 2ms/step -
 accuracy: 0.9862 - loss: 0.0430 - val_accuracy: 0.9024 - val_loss: 0.4681
 Epoch 85/128
 104/104 0s 2ms/step -
 accuracy: 0.9892 - loss: 0.0359 - val_accuracy: 0.8996 - val_loss: 0.4928
 Epoch 86/128
 104/104 0s 1ms/step -
 accuracy: 0.9909 - loss: 0.0328 - val_accuracy: 0.9031 - val_loss: 0.4879
 Epoch 87/128
 104/104 0s 1ms/step -
 accuracy: 0.9786 - loss: 0.0561 - val_accuracy: 0.9123 - val_loss: 0.4795
 Epoch 88/128
 104/104 0s 2ms/step -
 accuracy: 0.9896 - loss: 0.0370 - val_accuracy: 0.9066 - val_loss: 0.4756
 Epoch 89/128
 104/104 0s 2ms/step -
 accuracy: 0.9838 - loss: 0.0462 - val_accuracy: 0.9201 - val_loss: 0.4738
 Epoch 90/128
 104/104 0s 2ms/step -
 accuracy: 0.9903 - loss: 0.0320 - val_accuracy: 0.9088 - val_loss: 0.4950
 Epoch 91/128
 104/104 0s 2ms/step -
 accuracy: 0.9937 - loss: 0.0301 - val_accuracy: 0.9123 - val_loss: 0.4662
 Epoch 92/128
 104/104 0s 2ms/step -
 accuracy: 0.9910 - loss: 0.0305 - val_accuracy: 0.9109 - val_loss: 0.4935
 Epoch 93/128
 104/104 0s 2ms/step -
 accuracy: 0.9904 - loss: 0.0268 - val_accuracy: 0.9074 - val_loss: 0.4928
 Epoch 94/128
 104/104 0s 2ms/step -
 accuracy: 0.9728 - loss: 0.0724 - val_accuracy: 0.8975 - val_loss: 0.5042
 Epoch 95/128
 104/104 0s 2ms/step -
 accuracy: 0.9934 - loss: 0.0281 - val_accuracy: 0.9074 - val_loss: 0.4893
 Epoch 96/128
 104/104 0s 2ms/step -
 accuracy: 0.9920 - loss: 0.0290 - val_accuracy: 0.9130 - val_loss: 0.4947
 Epoch 97/128
 104/104 0s 2ms/step -

accuracy: 0.9948 - loss: 0.0234 - val_accuracy: 0.9059 - val_loss: 0.5050
 Epoch 98/128
 104/104 0s 2ms/step -
 accuracy: 0.9937 - loss: 0.0270 - val_accuracy: 0.9109 - val_loss: 0.4990
 Epoch 99/128
 104/104 0s 2ms/step -
 accuracy: 0.9891 - loss: 0.0347 - val_accuracy: 0.9052 - val_loss: 0.5158
 Epoch 100/128
 104/104 0s 2ms/step -
 accuracy: 0.9966 - loss: 0.0195 - val_accuracy: 0.9123 - val_loss: 0.5006
 Epoch 101/128
 104/104 0s 1ms/step -
 accuracy: 0.9925 - loss: 0.0245 - val_accuracy: 0.9088 - val_loss: 0.5225
 Epoch 102/128
 104/104 0s 1ms/step -
 accuracy: 0.9973 - loss: 0.0174 - val_accuracy: 0.9137 - val_loss: 0.5143
 Epoch 103/128
 104/104 0s 2ms/step -
 accuracy: 0.9905 - loss: 0.0283 - val_accuracy: 0.9010 - val_loss: 0.5611
 Epoch 104/128
 104/104 0s 2ms/step -
 accuracy: 0.9928 - loss: 0.0253 - val_accuracy: 0.9074 - val_loss: 0.5369
 Epoch 105/128
 104/104 0s 2ms/step -
 accuracy: 0.9926 - loss: 0.0288 - val_accuracy: 0.9038 - val_loss: 0.5575
 Epoch 106/128
 104/104 0s 2ms/step -
 accuracy: 0.9941 - loss: 0.0258 - val_accuracy: 0.9109 - val_loss: 0.5270
 Epoch 107/128
 104/104 0s 2ms/step -
 accuracy: 0.9913 - loss: 0.0334 - val_accuracy: 0.9017 - val_loss: 0.5566
 Epoch 108/128
 104/104 0s 2ms/step -
 accuracy: 0.9923 - loss: 0.0318 - val_accuracy: 0.9066 - val_loss: 0.5662
 Epoch 109/128
 104/104 0s 2ms/step -
 accuracy: 0.9863 - loss: 0.0321 - val_accuracy: 0.9066 - val_loss: 0.5528
 Epoch 110/128
 104/104 0s 2ms/step -
 accuracy: 0.9898 - loss: 0.0325 - val_accuracy: 0.9088 - val_loss: 0.5579
 Epoch 111/128
 104/104 0s 2ms/step -
 accuracy: 0.9956 - loss: 0.0155 - val_accuracy: 0.9123 - val_loss: 0.5588
 Epoch 112/128
 104/104 0s 2ms/step -
 accuracy: 0.9979 - loss: 0.0116 - val_accuracy: 0.9144 - val_loss: 0.5459
 Epoch 113/128
 104/104 0s 2ms/step -

accuracy: 0.9964 - loss: 0.0155 - val_accuracy: 0.9109 - val_loss: 0.5466
 Epoch 114/128
 104/104 0s 2ms/step -
 accuracy: 0.9974 - loss: 0.0137 - val_accuracy: 0.9123 - val_loss: 0.5539
 Epoch 115/128
 104/104 0s 2ms/step -
 accuracy: 0.9967 - loss: 0.0158 - val_accuracy: 0.9017 - val_loss: 0.5734
 Epoch 116/128
 104/104 0s 2ms/step -
 accuracy: 0.9845 - loss: 0.0471 - val_accuracy: 0.9059 - val_loss: 0.5916
 Epoch 117/128
 104/104 0s 2ms/step -
 accuracy: 0.9870 - loss: 0.0401 - val_accuracy: 0.9066 - val_loss: 0.6253
 Epoch 118/128
 104/104 0s 2ms/step -
 accuracy: 0.9830 - loss: 0.0431 - val_accuracy: 0.9088 - val_loss: 0.5822
 Epoch 119/128
 104/104 0s 2ms/step -
 accuracy: 0.9936 - loss: 0.0241 - val_accuracy: 0.9003 - val_loss: 0.6008
 Epoch 120/128
 104/104 0s 2ms/step -
 accuracy: 0.9930 - loss: 0.0189 - val_accuracy: 0.9024 - val_loss: 0.6023
 Epoch 121/128
 104/104 0s 2ms/step -
 accuracy: 0.9344 - loss: 0.3151 - val_accuracy: 0.8840 - val_loss: 0.5910
 Epoch 122/128
 104/104 0s 2ms/step -
 accuracy: 0.9806 - loss: 0.0628 - val_accuracy: 0.9038 - val_loss: 0.5987
 Epoch 123/128
 104/104 0s 2ms/step -
 accuracy: 0.9922 - loss: 0.0282 - val_accuracy: 0.9074 - val_loss: 0.5498
 Epoch 124/128
 104/104 0s 2ms/step -
 accuracy: 0.9988 - loss: 0.0145 - val_accuracy: 0.9116 - val_loss: 0.5673
 Epoch 125/128
 104/104 0s 2ms/step -
 accuracy: 0.9993 - loss: 0.0114 - val_accuracy: 0.9081 - val_loss: 0.5661
 Epoch 126/128
 104/104 0s 2ms/step -
 accuracy: 0.9991 - loss: 0.0109 - val_accuracy: 0.9095 - val_loss: 0.5727
 Epoch 127/128
 104/104 0s 2ms/step -
 accuracy: 0.9991 - loss: 0.0102 - val_accuracy: 0.9123 - val_loss: 0.5778
 Epoch 128/128
 104/104 0s 2ms/step -
 accuracy: 0.9989 - loss: 0.0118 - val_accuracy: 0.9052 - val_loss: 0.5981

```
[60]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Acurácia do modelo')
plt.ylabel('Acurácia')
plt.xlabel('Época')
plt.legend(['treinamento', 'validação'], loc='lower right')
plt.grid()
plt.show()
```



1.0.5 5. Apresente o resultado fazendo considerações sobre:

i. Como a complexidade do modelo impacta no desempenho? Um aumento da complexidade não parece fazer com que o modelo aumente sua acurácia no conjunto de validação, e ainda faz com que o *overfitting* aconteça mais rapidamente:

```
[61]: model_complex = Sequential()
model_complex.add(Dense(1024, input_dim = X_train.width, activation='relu')),
model_complex.add(Dense(512, activation='relu')),
model_complex.add(Dense(256, activation='relu')),
model_complex.add(Dense(128, activation='relu')),
```

```

model_complex.add(Dense(64, activation='relu')),
model_complex.add(Dense(32, activation='relu')),
model_complex.add(Dense(16, activation='relu')),
model_complex.add(Dense(Y_train.width, activation='softmax'))

optimizer = Adam()
model_complex.compile(optimizer=optimizer,
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])

history_complex = model_complex.fit(
    X_train.to_numpy(),
    Y_train.to_numpy(),
    epochs = 32,
    batch_size = 32,
    validation_data = (X_test.to_numpy(), Y_test.to_numpy()))

plt.plot(history_complex.history['accuracy'])
plt.plot(history_complex.history['val_accuracy'])
plt.title('Acurácia do modelo')
plt.ylabel('Acurácia')
plt.xlabel('Época')
plt.legend(['treinamento', 'validação'], loc='lower right')
plt.grid()
plt.show()

```

Epoch 1/32

```

/home/fbaltor/.pyenv/versions/3.11.10/lib/python3.11/site-
packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

104/104          2s 9ms/step -
accuracy: 0.5309 - loss: 1.2763 - val_accuracy: 0.7666 - val_loss: 0.6703

```

Epoch 2/32

```

104/104          1s 7ms/step -
accuracy: 0.8171 - loss: 0.5328 - val_accuracy: 0.8197 - val_loss: 0.5164

```

Epoch 3/32

```

104/104          1s 7ms/step -
accuracy: 0.8532 - loss: 0.4343 - val_accuracy: 0.8529 - val_loss: 0.4043

```

Epoch 4/32

```

104/104          1s 8ms/step -
accuracy: 0.8769 - loss: 0.3627 - val_accuracy: 0.8564 - val_loss: 0.4130

```

Epoch 5/32

```

104/104          1s 7ms/step -
accuracy: 0.8770 - loss: 0.3398 - val_accuracy: 0.8670 - val_loss: 0.3805

```

Epoch 6/32
104/104 1s 7ms/step -
accuracy: 0.8857 - loss: 0.3216 - val_accuracy: 0.8508 - val_loss: 0.4156
Epoch 7/32
104/104 1s 7ms/step -
accuracy: 0.8939 - loss: 0.2795 - val_accuracy: 0.8600 - val_loss: 0.4341
Epoch 8/32
104/104 1s 7ms/step -
accuracy: 0.8904 - loss: 0.2950 - val_accuracy: 0.8593 - val_loss: 0.4334
Epoch 9/32
104/104 1s 7ms/step -
accuracy: 0.9081 - loss: 0.2557 - val_accuracy: 0.8536 - val_loss: 0.4812
Epoch 10/32
104/104 1s 7ms/step -
accuracy: 0.8862 - loss: 0.2901 - val_accuracy: 0.8847 - val_loss: 0.3459
Epoch 11/32
104/104 1s 7ms/step -
accuracy: 0.9111 - loss: 0.2414 - val_accuracy: 0.8847 - val_loss: 0.3800
Epoch 12/32
104/104 1s 7ms/step -
accuracy: 0.9193 - loss: 0.2231 - val_accuracy: 0.8791 - val_loss: 0.4121
Epoch 13/32
104/104 1s 7ms/step -
accuracy: 0.9308 - loss: 0.2015 - val_accuracy: 0.8897 - val_loss: 0.3685
Epoch 14/32
104/104 1s 7ms/step -
accuracy: 0.9307 - loss: 0.1992 - val_accuracy: 0.8784 - val_loss: 0.3673
Epoch 15/32
104/104 1s 7ms/step -
accuracy: 0.9410 - loss: 0.1733 - val_accuracy: 0.8769 - val_loss: 0.4146
Epoch 16/32
104/104 1s 7ms/step -
accuracy: 0.9387 - loss: 0.1602 - val_accuracy: 0.8656 - val_loss: 0.4504
Epoch 17/32
104/104 1s 7ms/step -
accuracy: 0.9369 - loss: 0.1872 - val_accuracy: 0.8126 - val_loss: 0.6420
Epoch 18/32
104/104 1s 7ms/step -
accuracy: 0.8681 - loss: 0.4085 - val_accuracy: 0.8911 - val_loss: 0.3968
Epoch 19/32
104/104 1s 7ms/step -
accuracy: 0.9448 - loss: 0.1647 - val_accuracy: 0.8967 - val_loss: 0.3991
Epoch 20/32
104/104 1s 7ms/step -
accuracy: 0.9479 - loss: 0.1325 - val_accuracy: 0.8791 - val_loss: 0.4716
Epoch 21/32
104/104 1s 7ms/step -
accuracy: 0.9510 - loss: 0.1381 - val_accuracy: 0.8904 - val_loss: 0.3904

Epoch 22/32
104/104 1s 7ms/step -
accuracy: 0.9576 - loss: 0.1168 - val_accuracy: 0.8925 - val_loss: 0.3745

Epoch 23/32
104/104 1s 8ms/step -
accuracy: 0.9602 - loss: 0.1140 - val_accuracy: 0.9158 - val_loss: 0.3787

Epoch 24/32
104/104 1s 7ms/step -
accuracy: 0.9614 - loss: 0.1082 - val_accuracy: 0.8876 - val_loss: 0.5167

Epoch 25/32
104/104 1s 7ms/step -
accuracy: 0.9614 - loss: 0.1073 - val_accuracy: 0.9052 - val_loss: 0.4000

Epoch 26/32
104/104 1s 7ms/step -
accuracy: 0.9719 - loss: 0.0790 - val_accuracy: 0.9158 - val_loss: 0.4207

Epoch 27/32
104/104 1s 7ms/step -
accuracy: 0.9686 - loss: 0.0889 - val_accuracy: 0.9066 - val_loss: 0.3766


Epoch 28/32
104/104 1s 7ms/step -
accuracy: 0.9716 - loss: 0.0778 - val_accuracy: 0.9088 - val_loss: 0.4342

Epoch 29/32
104/104 1s 7ms/step -
accuracy: 0.9817 - loss: 0.0639 - val_accuracy: 0.9003 - val_loss: 0.4609

Epoch 30/32
104/104 1s 7ms/step -
accuracy: 0.9792 - loss: 0.0620 - val_accuracy: 0.9144 - val_loss: 0.4417

Epoch 31/32
104/104 1s 7ms/step -
accuracy: 0.9841 - loss: 0.0574 - val_accuracy: 0.9024 - val_loss: 0.4987

Epoch 32/32
104/104 1s 7ms/step -
accuracy: 0.9849 - loss: 0.0462 - val_accuracy: 0.9194 - val_loss: 0.4734



ii. Comente sobre a estratégia de ajuste dos hiper parâmetros adotada para ajustar o modelo.

- O *learning rate* foi deixado em seu valor padrão.
- A complexidade (número de camadas e neurônios de cada camada) foi deixada baixa porque o modelo facilmente apresenta *overfitting*.
- O número de épocas não precisou ser alto (32) dado que rapidamente o treinamento atinge um bom valor de acurácia, ao passo que valores maiores também causam *overfitting*.
- Foram testadas algumas outras funções de ativação (*tanh*, *softmax* em todas as camadas, etc), mas elas não pareceram ser tão determinantes no resultado do treinamento.

iii. Como a quantidade de amostras de cada classe pode influenciar no desempenho da rede neural? A quantidade de amostras de cada classe e seu balanceamento em especial se mostrou extremamente importante para o bom desempenho da rede neural. Abaixo segue o treinamento do modelo com os dados desbalanceados:

```
[62]: train_fraction = 0.7
X_train, X_test, Y_train, Y_test = train_test_split(X_unbalanced, Y_unbalanced,
↳ test_size = (1 - train_fraction), shuffle = True)
```

```

model_unb = Sequential()
model_unb.add(Dense(200, activation='relu')),
model_unb.add(Dense(50, activation='relu')),
model_unb.add(Dense(10, activation='relu')),
model_unb.add(Dense(Y_train.width, activation='softmax'))

optimizer = Adam()
model_unb.compile(optimizer=optimizer,
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history_unb = model_unb.fit(
    X_train.to_numpy(),
    Y_train.to_numpy(),
    epochs = 128,
    batch_size = 32,
    validation_data = (X_test.to_numpy(), Y_test.to_numpy()))

plt.plot(history_unb.history['accuracy'])
plt.plot(history_unb.history['val_accuracy'])
plt.title('Acurácia do modelo')
plt.ylabel('Acurácia')
plt.xlabel('Época')
plt.legend(['treinamento', 'validação'], loc='lower right')
plt.grid()
plt.show()

```

Epoch 1/128

43/43 1s 5ms/step -

accuracy: 0.3226 - loss: 1.7882 - val_accuracy: 0.5506 - val_loss: 1.3424

Epoch 2/128

43/43 0s 2ms/step -

accuracy: 0.5232 - loss: 1.3009 - val_accuracy: 0.6261 - val_loss: 1.0726

Epoch 3/128

43/43 0s 2ms/step -

accuracy: 0.6419 - loss: 1.0167 - val_accuracy: 0.6587 - val_loss: 0.9640

Epoch 4/128

43/43 0s 2ms/step -

accuracy: 0.6427 - loss: 0.9565 - val_accuracy: 0.6690 - val_loss: 0.9017

Epoch 5/128

43/43 0s 2ms/step -

accuracy: 0.6946 - loss: 0.8288 - val_accuracy: 0.6844 - val_loss: 0.8661

Epoch 6/128

43/43 0s 2ms/step -

accuracy: 0.6963 - loss: 0.7920 - val_accuracy: 0.6913 - val_loss: 0.8385

Epoch 7/128

43/43 0s 2ms/step -

accuracy: 0.7041 - loss: 0.7618 - val_accuracy: 0.6913 - val_loss: 0.8207
 Epoch 8/128
 43/43 0s 2ms/step -
 accuracy: 0.7213 - loss: 0.7384 - val_accuracy: 0.7067 - val_loss: 0.7921
 Epoch 9/128
 43/43 0s 2ms/step -
 accuracy: 0.7479 - loss: 0.6574 - val_accuracy: 0.7204 - val_loss: 0.7859
 Epoch 10/128
 43/43 0s 2ms/step -
 accuracy: 0.7536 - loss: 0.6755 - val_accuracy: 0.6913 - val_loss: 0.7842
 Epoch 11/128
 43/43 0s 2ms/step -
 accuracy: 0.7705 - loss: 0.6409 - val_accuracy: 0.7376 - val_loss: 0.7469
 Epoch 12/128
 43/43 0s 2ms/step -
 accuracy: 0.7760 - loss: 0.6132 - val_accuracy: 0.7358 - val_loss: 0.7153
 Epoch 13/128
 43/43 0s 2ms/step -
 accuracy: 0.7876 - loss: 0.5843 - val_accuracy: 0.7256 - val_loss: 0.7448
 Epoch 14/128
 43/43 0s 2ms/step -
 accuracy: 0.7875 - loss: 0.5765 - val_accuracy: 0.7256 - val_loss: 0.7560
 Epoch 15/128
 43/43 0s 2ms/step -
 accuracy: 0.8031 - loss: 0.5477 - val_accuracy: 0.7238 - val_loss: 0.7247
 Epoch 16/128
 43/43 0s 2ms/step -
 accuracy: 0.7936 - loss: 0.5488 - val_accuracy: 0.7427 - val_loss: 0.7031
 Epoch 17/128
 43/43 0s 2ms/step -
 accuracy: 0.7918 - loss: 0.5589 - val_accuracy: 0.7324 - val_loss: 0.6958
 Epoch 18/128
 43/43 0s 2ms/step -
 accuracy: 0.8003 - loss: 0.5130 - val_accuracy: 0.7033 - val_loss: 0.7492
 Epoch 19/128
 43/43 0s 2ms/step -
 accuracy: 0.8337 - loss: 0.4723 - val_accuracy: 0.7479 - val_loss: 0.6996
 Epoch 20/128
 43/43 0s 3ms/step -
 accuracy: 0.8174 - loss: 0.4755 - val_accuracy: 0.7273 - val_loss: 0.6981
 Epoch 21/128
 43/43 0s 2ms/step -
 accuracy: 0.8270 - loss: 0.4859 - val_accuracy: 0.7376 - val_loss: 0.6983
 Epoch 22/128
 43/43 0s 2ms/step -
 accuracy: 0.8240 - loss: 0.4509 - val_accuracy: 0.7496 - val_loss: 0.6958
 Epoch 23/128
 43/43 0s 2ms/step -

accuracy: 0.8417 - loss: 0.4484 - val_accuracy: 0.7479 - val_loss: 0.7122
 Epoch 24/128
 43/43 0s 2ms/step -
 accuracy: 0.8457 - loss: 0.4184 - val_accuracy: 0.7290 - val_loss: 0.7119
 Epoch 25/128
 43/43 0s 2ms/step -
 accuracy: 0.8371 - loss: 0.4315 - val_accuracy: 0.7221 - val_loss: 0.7407
 Epoch 26/128
 43/43 0s 2ms/step -
 accuracy: 0.8564 - loss: 0.4162 - val_accuracy: 0.7358 - val_loss: 0.7151
 Epoch 27/128
 43/43 0s 2ms/step -
 accuracy: 0.8447 - loss: 0.4110 - val_accuracy: 0.7427 - val_loss: 0.7057
 Epoch 28/128
 43/43 0s 2ms/step -
 accuracy: 0.8606 - loss: 0.3798 - val_accuracy: 0.7513 - val_loss: 0.7042
 Epoch 29/128
 43/43 0s 2ms/step -
 accuracy: 0.8582 - loss: 0.3933 - val_accuracy: 0.7581 - val_loss: 0.6941
 Epoch 30/128
 43/43 0s 2ms/step -
 accuracy: 0.8645 - loss: 0.3733 - val_accuracy: 0.7513 - val_loss: 0.7069
 Epoch 31/128
 43/43 0s 2ms/step -
 accuracy: 0.8625 - loss: 0.3900 - val_accuracy: 0.7461 - val_loss: 0.7105
 Epoch 32/128
 43/43 0s 2ms/step -
 accuracy: 0.8764 - loss: 0.3440 - val_accuracy: 0.7015 - val_loss: 0.8082
 Epoch 33/128
 43/43 0s 2ms/step -
 accuracy: 0.8480 - loss: 0.4086 - val_accuracy: 0.7341 - val_loss: 0.7143
 Epoch 34/128
 43/43 0s 2ms/step -
 accuracy: 0.8759 - loss: 0.3379 - val_accuracy: 0.7513 - val_loss: 0.7258
 Epoch 35/128
 43/43 0s 2ms/step -
 accuracy: 0.8581 - loss: 0.3672 - val_accuracy: 0.7393 - val_loss: 0.7283
 Epoch 36/128
 43/43 0s 2ms/step -
 accuracy: 0.8804 - loss: 0.3286 - val_accuracy: 0.7410 - val_loss: 0.7216
 Epoch 37/128
 43/43 0s 2ms/step -
 accuracy: 0.8602 - loss: 0.3475 - val_accuracy: 0.7410 - val_loss: 0.7784
 Epoch 38/128
 43/43 0s 2ms/step -
 accuracy: 0.8815 - loss: 0.3340 - val_accuracy: 0.7564 - val_loss: 0.7284
 Epoch 39/128
 43/43 0s 2ms/step -

accuracy: 0.8718 - loss: 0.3353 - val_accuracy: 0.7427 - val_loss: 0.7664
 Epoch 40/128
 43/43 0s 2ms/step -
 accuracy: 0.8805 - loss: 0.3079 - val_accuracy: 0.7256 - val_loss: 0.8076
 Epoch 41/128
 43/43 0s 2ms/step -
 accuracy: 0.8810 - loss: 0.3140 - val_accuracy: 0.7410 - val_loss: 0.7567
 Epoch 42/128
 43/43 0s 2ms/step -
 accuracy: 0.9043 - loss: 0.2904 - val_accuracy: 0.7444 - val_loss: 0.7748
 Epoch 43/128
 43/43 0s 2ms/step -
 accuracy: 0.8950 - loss: 0.3011 - val_accuracy: 0.7256 - val_loss: 0.8031
 Epoch 44/128
 43/43 0s 2ms/step -
 accuracy: 0.8770 - loss: 0.3447 - val_accuracy: 0.7479 - val_loss: 0.7744
 Epoch 45/128
 43/43 0s 2ms/step -
 accuracy: 0.8907 - loss: 0.2898 - val_accuracy: 0.7376 - val_loss: 0.8016
 Epoch 46/128
 43/43 0s 2ms/step -
 accuracy: 0.9043 - loss: 0.2797 - val_accuracy: 0.7513 - val_loss: 0.7982
 Epoch 47/128
 43/43 0s 2ms/step -
 accuracy: 0.9076 - loss: 0.2771 - val_accuracy: 0.7444 - val_loss: 0.7923
 Epoch 48/128
 43/43 0s 2ms/step -
 accuracy: 0.9239 - loss: 0.2418 - val_accuracy: 0.7324 - val_loss: 0.8137
 Epoch 49/128
 43/43 0s 2ms/step -
 accuracy: 0.9198 - loss: 0.2405 - val_accuracy: 0.7307 - val_loss: 0.8159
 Epoch 50/128
 43/43 0s 2ms/step -
 accuracy: 0.9244 - loss: 0.2352 - val_accuracy: 0.7290 - val_loss: 0.7956
 Epoch 51/128
 43/43 0s 2ms/step -
 accuracy: 0.9078 - loss: 0.2505 - val_accuracy: 0.7599 - val_loss: 0.8369
 Epoch 52/128
 43/43 0s 2ms/step -
 accuracy: 0.9171 - loss: 0.2580 - val_accuracy: 0.7307 - val_loss: 0.8638
 Epoch 53/128
 43/43 0s 2ms/step -
 accuracy: 0.9068 - loss: 0.2606 - val_accuracy: 0.7410 - val_loss: 0.8481
 Epoch 54/128
 43/43 0s 2ms/step -
 accuracy: 0.9106 - loss: 0.2371 - val_accuracy: 0.7393 - val_loss: 0.8461
 Epoch 55/128
 43/43 0s 3ms/step -


accuracy: 0.9257 - loss: 0.2253 - val_accuracy: 0.7238 - val_loss: 0.8672
 Epoch 56/128
 43/43 0s 2ms/step -
 accuracy: 0.9115 - loss: 0.2616 - val_accuracy: 0.7324 - val_loss: 0.8661
 Epoch 57/128
 43/43 0s 2ms/step -
 accuracy: 0.9410 - loss: 0.2011 - val_accuracy: 0.7444 - val_loss: 0.8474
 Epoch 58/128
 43/43 0s 2ms/step -
 accuracy: 0.9233 - loss: 0.2323 - val_accuracy: 0.7324 - val_loss: 0.8842
 Epoch 59/128
 43/43 0s 2ms/step -
 accuracy: 0.9268 - loss: 0.2017 - val_accuracy: 0.7616 - val_loss: 0.8517
 Epoch 60/128
 43/43 0s 2ms/step -
 accuracy: 0.9364 - loss: 0.1942 - val_accuracy: 0.7513 - val_loss: 0.8841
 Epoch 61/128
 43/43 0s 2ms/step -
 accuracy: 0.9366 - loss: 0.1983 - val_accuracy: 0.7496 - val_loss: 0.9224
 Epoch 62/128
 43/43 0s 2ms/step -
 accuracy: 0.9383 - loss: 0.1985 - val_accuracy: 0.7393 - val_loss: 0.9013
 Epoch 63/128
 43/43 0s 2ms/step -
 accuracy: 0.9318 - loss: 0.2033 - val_accuracy: 0.7376 - val_loss: 0.9106
 Epoch 64/128
 43/43 0s 2ms/step -
 accuracy: 0.9434 - loss: 0.1806 - val_accuracy: 0.7513 - val_loss: 0.9144
 Epoch 65/128
 43/43 0s 2ms/step -
 accuracy: 0.9529 - loss: 0.1641 - val_accuracy: 0.7290 - val_loss: 0.9320
 Epoch 66/128
 43/43 0s 2ms/step -
 accuracy: 0.9509 - loss: 0.1689 - val_accuracy: 0.7358 - val_loss: 0.9360
 Epoch 67/128
 43/43 0s 2ms/step -
 accuracy: 0.9509 - loss: 0.1694 - val_accuracy: 0.7461 - val_loss: 0.9404
 Epoch 68/128
 43/43 0s 2ms/step -
 accuracy: 0.9312 - loss: 0.1803 - val_accuracy: 0.7341 - val_loss: 0.9368
 Epoch 69/128
 43/43 0s 2ms/step -
 accuracy: 0.9421 - loss: 0.1722 - val_accuracy: 0.7324 - val_loss: 0.9790
 Epoch 70/128
 43/43 0s 2ms/step -
 accuracy: 0.9537 - loss: 0.1615 - val_accuracy: 0.7427 - val_loss: 0.9615
 Epoch 71/128
 43/43 0s 2ms/step -

accuracy: 0.9590 - loss: 0.1398 - val_accuracy: 0.7616 - val_loss: 0.9311
 Epoch 72/128
 43/43 0s 2ms/step -
 accuracy: 0.9480 - loss: 0.1493 - val_accuracy: 0.7427 - val_loss: 0.9914
 Epoch 73/128
 43/43 0s 2ms/step -
 accuracy: 0.9520 - loss: 0.1637 - val_accuracy: 0.7461 - val_loss: 0.9855
 Epoch 74/128
 43/43 0s 2ms/step -
 accuracy: 0.9607 - loss: 0.1377 - val_accuracy: 0.7410 - val_loss: 0.9964
 Epoch 75/128
 43/43 0s 2ms/step -
 accuracy: 0.9596 - loss: 0.1381 - val_accuracy: 0.7393 - val_loss: 1.0230
 Epoch 76/128
 43/43 0s 2ms/step -
 accuracy: 0.9490 - loss: 0.1516 - val_accuracy: 0.7461 - val_loss: 1.0385
 Epoch 77/128
 43/43 0s 3ms/step -
 accuracy: 0.9506 - loss: 0.1492 - val_accuracy: 0.7290 - val_loss: 1.0482
 Epoch 78/128
 43/43 0s 4ms/step -
 accuracy: 0.9447 - loss: 0.1539 - val_accuracy: 0.7427 - val_loss: 1.0549
 Epoch 79/128
 43/43 0s 3ms/step -
 accuracy: 0.9702 - loss: 0.1251 - val_accuracy: 0.7530 - val_loss: 1.0325
 Epoch 80/128
 43/43 0s 3ms/step -
 accuracy: 0.9618 - loss: 0.1307 - val_accuracy: 0.7513 - val_loss: 1.0254
 Epoch 81/128
 43/43 0s 3ms/step -
 accuracy: 0.9576 - loss: 0.1278 - val_accuracy: 0.7256 - val_loss: 1.1121
 Epoch 82/128
 43/43 0s 3ms/step -
 accuracy: 0.9570 - loss: 0.1370 - val_accuracy: 0.7599 - val_loss: 1.0917
 Epoch 83/128
 43/43 0s 3ms/step -
 accuracy: 0.9621 - loss: 0.1151 - val_accuracy: 0.7461 - val_loss: 1.0856
 Epoch 84/128
 43/43 0s 3ms/step -
 accuracy: 0.9512 - loss: 0.1389 - val_accuracy: 0.7410 - val_loss: 1.0946
 Epoch 85/128
 43/43 0s 3ms/step -
 accuracy: 0.9654 - loss: 0.1136 - val_accuracy: 0.7358 - val_loss: 1.1389
 Epoch 86/128
 43/43 0s 3ms/step -
 accuracy: 0.9619 - loss: 0.1141 - val_accuracy: 0.7410 - val_loss: 1.1435
 Epoch 87/128
 43/43 0s 3ms/step -

accuracy: 0.9678 - loss: 0.1005 - val_accuracy: 0.7307 - val_loss: 1.1410
 Epoch 88/128
 43/43 0s 3ms/step -
 accuracy: 0.9700 - loss: 0.1087 - val_accuracy: 0.7530 - val_loss: 1.1279
 Epoch 89/128
 43/43 0s 2ms/step -
 accuracy: 0.9735 - loss: 0.0988 - val_accuracy: 0.7461 - val_loss: 1.1329
 Epoch 90/128
 43/43 0s 2ms/step -
 accuracy: 0.9647 - loss: 0.1011 - val_accuracy: 0.7376 - val_loss: 1.1983
 Epoch 91/128
 43/43 0s 3ms/step -
 accuracy: 0.9689 - loss: 0.1002 - val_accuracy: 0.7410 - val_loss: 1.1373
 Epoch 92/128
 43/43 0s 3ms/step -
 accuracy: 0.9742 - loss: 0.0964 - val_accuracy: 0.7479 - val_loss: 1.1706
 Epoch 93/128
 43/43 0s 4ms/step -
 accuracy: 0.9769 - loss: 0.0867 - val_accuracy: 0.7410 - val_loss: 1.1779
 Epoch 94/128
 43/43 0s 3ms/step -
 accuracy: 0.9670 - loss: 0.1152 - val_accuracy: 0.7479 - val_loss: 1.1886
 Epoch 95/128
 43/43 0s 3ms/step -
 accuracy: 0.9737 - loss: 0.0970 - val_accuracy: 0.7581 - val_loss: 1.1832
 Epoch 96/128
 43/43 0s 3ms/step -
 accuracy: 0.9746 - loss: 0.1042 - val_accuracy: 0.7376 - val_loss: 1.2407
 Epoch 97/128
 43/43 0s 3ms/step -
 accuracy: 0.9771 - loss: 0.0922 - val_accuracy: 0.7324 - val_loss: 1.2466
 Epoch 98/128
 43/43 0s 3ms/step -
 accuracy: 0.9816 - loss: 0.0884 - val_accuracy: 0.7461 - val_loss: 1.2517
 Epoch 99/128
 43/43 0s 3ms/step -
 accuracy: 0.9734 - loss: 0.0872 - val_accuracy: 0.7324 - val_loss: 1.3254
 Epoch 100/128
 43/43 0s 3ms/step -
 accuracy: 0.9567 - loss: 0.1140 - val_accuracy: 0.7307 - val_loss: 1.2896
 Epoch 101/128
 43/43 0s 2ms/step -
 accuracy: 0.9772 - loss: 0.0891 - val_accuracy: 0.7393 - val_loss: 1.2625
 Epoch 102/128
 43/43 0s 2ms/step -
 accuracy: 0.9798 - loss: 0.0937 - val_accuracy: 0.7256 - val_loss: 1.3064
 Epoch 103/128
 43/43 0s 3ms/step -

accuracy: 0.9780 - loss: 0.0795 - val_accuracy: 0.7376 - val_loss: 1.2794
 Epoch 104/128
 43/43 0s 2ms/step -
 accuracy: 0.9770 - loss: 0.0725 - val_accuracy: 0.7479 - val_loss: 1.2851
 Epoch 105/128
 43/43 0s 3ms/step -
 accuracy: 0.9798 - loss: 0.0774 - val_accuracy: 0.7376 - val_loss: 1.2928
 Epoch 106/128
 43/43 0s 3ms/step -
 accuracy: 0.9845 - loss: 0.0670 - val_accuracy: 0.7461 - val_loss: 1.3410
 Epoch 107/128
 43/43 0s 5ms/step -
 accuracy: 0.9789 - loss: 0.0694 - val_accuracy: 0.7376 - val_loss: 1.3023
 Epoch 108/128
 43/43 0s 3ms/step -
 accuracy: 0.9865 - loss: 0.0644 - val_accuracy: 0.7479 - val_loss: 1.3215
 Epoch 109/128
 43/43 0s 3ms/step -
 accuracy: 0.9856 - loss: 0.0689 - val_accuracy: 0.7479 - val_loss: 1.3904
 Epoch 110/128
 43/43 0s 3ms/step -
 accuracy: 0.9723 - loss: 0.0791 - val_accuracy: 0.7530 - val_loss: 1.3569
 Epoch 111/128
 43/43 0s 3ms/step -
 accuracy: 0.9794 - loss: 0.0719 - val_accuracy: 0.7256 - val_loss: 1.4651
 Epoch 112/128
 43/43 0s 2ms/step -
 accuracy: 0.9773 - loss: 0.0805 - val_accuracy: 0.7358 - val_loss: 1.4272
 Epoch 113/128
 43/43 0s 3ms/step -
 accuracy: 0.9905 - loss: 0.0668 - val_accuracy: 0.7273 - val_loss: 1.4057
 Epoch 114/128
 43/43 0s 3ms/step -
 accuracy: 0.9777 - loss: 0.0673 - val_accuracy: 0.7273 - val_loss: 1.4098
 Epoch 115/128
 43/43 0s 3ms/step -
 accuracy: 0.9892 - loss: 0.0573 - val_accuracy: 0.7341 - val_loss: 1.5103
 Epoch 116/128
 43/43 0s 3ms/step -
 accuracy: 0.9701 - loss: 0.0840 - val_accuracy: 0.7427 - val_loss: 1.4440
 Epoch 117/128
 43/43 0s 3ms/step -
 accuracy: 0.9813 - loss: 0.0741 - val_accuracy: 0.7461 - val_loss: 1.3822
 Epoch 118/128
 43/43 0s 3ms/step -
 accuracy: 0.9887 - loss: 0.0563 - val_accuracy: 0.7427 - val_loss: 1.4268
 Epoch 119/128
 43/43 0s 4ms/step -

accuracy: 0.9920 - loss: 0.0549 - val_accuracy: 0.7273 - val_loss: 1.4671
Epoch 120/128
43/43 0s 3ms/step -
accuracy: 0.9904 - loss: 0.0500 - val_accuracy: 0.7393 - val_loss: 1.4386
Epoch 121/128
43/43 0s 3ms/step -
accuracy: 0.9899 - loss: 0.0521 - val_accuracy: 0.7461 - val_loss: 1.4628
Epoch 122/128
43/43 0s 3ms/step -
accuracy: 0.9842 - loss: 0.0579 - val_accuracy: 0.7341 - val_loss: 1.5026
Epoch 123/128
43/43 0s 3ms/step -
accuracy: 0.9896 - loss: 0.0540 - val_accuracy: 0.7341 - val_loss: 1.5182
Epoch 124/128
43/43 0s 3ms/step -
accuracy: 0.9852 - loss: 0.0524 - val_accuracy: 0.7324 - val_loss: 1.5532
Epoch 125/128
43/43 0s 3ms/step -
accuracy: 0.9857 - loss: 0.0575 - val_accuracy: 0.7444 - val_loss: 1.4930
Epoch 126/128
43/43 0s 3ms/step -
accuracy: 0.9899 - loss: 0.0466 - val_accuracy: 0.7427 - val_loss: 1.5143
Epoch 127/128
43/43 0s 3ms/step -
accuracy: 0.9902 - loss: 0.0444 - val_accuracy: 0.7461 - val_loss: 1.4919
Epoch 128/128
43/43 0s 3ms/step -
accuracy: 0.9856 - loss: 0.0496 - val_accuracy: 0.7376 - val_loss: 1.5692



A piora do seu desempenho é nítida, com o *overfitting* ainda mais presente e a acurácia estagnando quase 15% abaixo do melhor modelo obtido anteriormente.

iv. Quais variáveis de entrada são mais relevantes para o problema? Vamos empregar uma técnica simples, obtida com ajuda do [Gepeto](#) e que consiste em calcular uma acurácia básica (*baseline*) com o modelo completo para então randomizar os valores de cada *feature* individualmente e calcular essa acurácia específica:

```
[72]: from sklearn.metrics import accuracy_score

def permutation_importance(model, X_val, y_val):
    baseline = accuracy_score(y_val, model.predict(X_val).round())
    importances = []

    for col in range(X_val.shape[1]):
        X_temp = X_val.copy()
        np.random.shuffle(X_temp[:, col]) # Shuffle the feature column
        shuffled = accuracy_score(y_val, model.predict(X_temp).round())
        importances.append(baseline - shuffled) # Calculate importance
```



```
return importances
```

```
importances = permutation_importance(model, X.to_numpy(), Y.to_numpy())
```

```
148/148      0s 601us/step
148/148      0s 575us/step
148/148      0s 596us/step
148/148      0s 635us/step
148/148      0s 626us/step
148/148      0s 704us/step
148/148      0s 628us/step
148/148      0s 625us/step
148/148      0s 675us/step
148/148      0s 670us/step
148/148      0s 681us/step
148/148      0s 650us/step
148/148      0s 606us/step
148/148      0s 622us/step
148/148      0s 601us/step
148/148      0s 628us/step
148/148      0s 579us/step
148/148      0s 640us/step
148/148      0s 611us/step
148/148      0s 606us/step
148/148      0s 643us/step
148/148      0s 590us/step
148/148      0s 586us/step
148/148      0s 579us/step
```

```
[88]: feature_importance = pl.DataFrame(
      {
          "feature": X.columns,
          "importance": importances
      }
  )

  feature_importance = feature_importance.sort(by = "importance").reverse()

  with pl.Config(tbl_rows = -1):
      print(feature_importance)
```

```
shape: (23, 2)
```

feature	importance
---	---
str	f64
Length_of_Conveyer	0.184037

Square_Index	0.165145
Log_Y_Index	0.157504
X_Minimum	0.132668
Edges_X_Index	0.130758
Maximum_of_Luminosity	0.1263
Y_Minimum	0.121842
SigmoidOfAreas	0.118022
Steel_Plate_Thickness	0.114201
Edges_Index	0.097644
Minimum_of_Luminosity	0.09637
TypeOfSteel_A300	0.091063
Log_X_Index	0.08703
Edges_Y_Index	0.081724
Empty_Index	0.077054
Orientation_Index	0.068987
Outside_Global_Index	0.065591
Luminosity_Index	0.063893
TypeOfSteel_A400	0.055827
LogOfAreas	0.050732
Outside_X_Index	0.021651
Y_Perimeter	0.010613
Pixels_Areas	0.009552

Percebemos que as variáveis de entrada mais importantes são `Length_of_Conveyer`, `Square_Index` e `Log_Y_Index`. Entretanto, a importância das variáveis parece decrescer gradualmente, de maneira que é difícil dizer que há um conjunto muito mais determinante que os demais.